

A Clinical-Oriented AI Framework for Multi-Class Kidney Abnormality

Detection and Segmentation

A Thesis Project

Presented to the

School of Information Technology

Mapua University

In Partial Fulfillment

of the Requirements for the Degree

Bachelor of Science in Computer Science

by

Yñikko Arzee Neo D. Aguas

Carl Emmanuel M. Macabales

Robin Jairic T. Macatangay

February 2026



CAPSTONE / THESIS PROPOSAL APPROVAL FORM

Document No: FM-SIT-08-01

Effective Date: Aug. 22, 2017

This capstone / thesis proposal paper entitled :

A Clinical-Oriented AI Framework for Multi-Class Kidney Abnormality
Detection and Segmentation


prepared and presented by the following students on February 27, 2026

Yñikko Arzee Neo D. Aguas

Carl Emmanuel M. Macabales

Robin Jairic T. Macatangay


was supervised and endorsed by the advisor/s


Dr. Lysa V. Comia

Name & Signature

Name & Signature


and hereby approved and recommended to be accepted
as fulfillment to the requirement of the course Computer Science by:


Joseph Albert E. Mazo

Name & Signature


Dr. Richard N. Monreal

Name & Signature


Elijah Erle O. Reyes

Name & Signature

Noted by:



Dr. John Paul Q. Tomas

Name & Signature of Course Coordinator

Date

RECOMMENDATION SHEET

This is to certify that I have supervised the preparation of and read the thesis proposal prepared by **Yñikko Arzee Neo D. Aguas, Carl Emmanuel M. Macabales, and Robin Jairic T. Macatangay** entitled **A Clinical-Oriented AI Framework for Multi-Class Kidney Abnormality Detection and Segmentation** and that the said paper has been submitted for final examination by the Oral Examination Committee.



Dr. Lysa V. Comia
Thesis Adviser

TABLE OF CONTENTS

TITLE PAGE	i
RECOMMENDATION SHEET.....	ii
TABLE OF CONTENTS.....	iii
LIST OF FIGURES	viii
LIST OF TABLES	ix
Chapter 1	1
1.1. Context of the Study	1
1.2. The Problem, Gap or Opportunity	6
1.3. Research Questions.....	8
1.4. Research Objectives.....	9
1.5 Significance of the Study	10
1.6 Scope and Limitations (or Delimitations).....	12
Chapter 2.....	14
2.1 Clinical Significance of Kidney Abnormalities.....	14
2.1.1 Imaging Characteristics of Renal Abnormalities.....	14
2.1.2 Imaging Characterization Systems and Classification.....	15
2.2 Object Detection Approaches in Urology.....	16
2.2.1 Object Detection in Medical Imaging: Concepts and Approaches.....	16
2.2.2 Detection vs. Segmentation: Complementary Methodologies.....	17

2.2.3 Region-Based Deep Learning Detection Methods	19
2.2.4 Single-Shot and Real-Time Detection Architectures.....	21
2.2.5 Transformer and Advanced Detection Models	23
2.2.6 Detection and Segmentation Integration in Urological Clinical Workflows...	25
2.3 Evolution of U-Net in Medical Image Segmentation	26
2.3.1 What is U-Net	26
2.3.2 Architectural Components and Design Philosophy	28
2.3.3 Clinical Applications and Modality Coverage.....	29
2.3.4 U-Net Variants and Architectural Innovations	30
2.3.5 Performance Metrics and Clinical Benchmarking.....	33
2.4 The nnU-Net Framework.....	37
2.4.1 Overview of nnU-Net	37
2.4.2 Automated Configuration Architecture	38
2.4.3 Cascade Architecture and Multi-Resolution Processing.....	39
2.4.4 Cross-Institutional Robustness and Domain Generalization	40
2.4.5 Empirical Performance and Validation.....	42
2.4.6 Clinical Translation and Accessibility	42
2.4.7 State-of-the-Art Deep Learning Frameworks for Kidney Segmentation.....	43
2.5 Identified Research Gaps	44
2.5.1 Single-Class Focus in Existing Segmentation Systems	44

2.5.2 Fragmented Public Datasets and the Need for Data Unification	45
2.5.3 Insufficient Integration of Clinical Classification Systems	47
2.5.4 The Comprehensive Metric Evaluation Gap.....	48
2.5.5 Gap Summary and Research Opportunity	49
Chapter 3	50
3.1 Conceptual Framework.....	50
3.1.1 The Input-Process-Output (IPO) Model	51
3.1.2 Input Phase: Data Ingestion and Physiological Standardization.....	52
3.1.3 Process Phase: Deep Learning and Volumetric Segmentation	53
3.1.4 Output Phase: Evaluation and Clinical Inference	53
3.2 Data Management	54
3.2.1 Sources and Ethical Considerations.....	55
3.2.2 Ethical Compliance Framework:	57
3.2.3 DICOM to NIfTI Conversion	58
3.2.4 Dataset Organization and Structure	59
3.2.5 Dataset Characterization	60
3.3 Data Preprocessing.....	61
3.3.1 Hounsfield Unit (HU) Clipping and Intensity Filtering.....	62
3.3.2 Physiological Justification for Clipping Range:	63
3.3.3 Isotropic Resampling and Spatial Normalization	64

3.3.4 Cropping to Non-Zero Region	66
3.3.5 Augmentation Setting and Data Splitting	66
3.4 Network Architecture: The 3D U-Net and Deep Supervision	68
3.4.1 Architectural Topology	69
3.4.2 Deep Supervision	71
3.5 Computational Pipeline.....	72
3.5.1 Environment and Reproducibility	72
3.5.2 Training Configuration and Optimization	73
3.5.3 Loss Function Math	73
3.5.4 Inference and Post-Processing	74
3.6 Annotation Pipeline	75
3.6.1 Domain Expert Profile	76
3.7 Evaluation Metrics	79
3.7.1 Dice Similarity Coefficient (DSC).....	79
3.7.2 Intersection over Union (IoU).....	79
3.7.3 95th Percentile Hausdorff Distance (HD95).....	80
3.7.4 Precision, Recall, and F1-Score	81
3.7.5 Per-Class Analysis	81
3.7.6 Computational Efficiency Metrics.....	82
3.7.7 Baseline Benchmarks.....	82

3.8 Failure Analysis and Clinical Interpretability	84
3.9 Clinical Feature Quantification.....	85
3.9.1 Volumetric Calculation.....	85
3.9.2 Maximum Diameter (Tumor Size).....	85
3.10 Testing and Validation Pipeline.....	86
3.10.1 Experimental Configuration Matrix.....	86
3.10.2 Held-Out Test Set Preparation	87
3.10.3 Inference & Post-Processing.....	87
3.10.4 Metric Computation	88
3.10.5 Computational Efficiency Recording	89
3.10.6 Cross Configuration Comparison	89
3.10.7 Domain Expert Review.....	90
3.11 Web Deployment	90
3.11.1 Framework Selection: Streamlit	90
3.11.2 Application Architecture.....	91
3.11.3 Clinical Reporting Features	92
REFERENCES	93

LIST OF FIGURES

<i>Figure 1. Comparison of Object Detection versus Semantic Segmentation in Medical Imaging. Adapted from [28].</i>	18
<i>Figure 2. Faster R-CNN Architecture. Reprinted from [15].</i>	20
<i>Figure 3. YOLO Architecture. Reprinted from [34].</i>	21
<i>Figure 4. Vision Transformers. Reprinted from [38].</i>	24
<i>Figure 5. The U-Net Architecture. Illustration of the contracting and expanding paths with skip connections. Reprinted from [12].</i>	27
<i>Figure 6. Modality Versatility.</i>	30
<i>Figure 7. U-Net++ Architecture. Reprinted from [42].</i>	31
<i>Figure 8. Attention Mechanism. Reprinted from [43].</i>	32
<i>Figure 9. Hybrid Transformers. Reprinted from [46].</i>	33
<i>Figure 10. Volumetric Overlap Metrics.</i>	34
<i>Figure 11. Boundary Distance Metrics. Reprinted from [50].</i>	36
<i>Figure 12. nnU-Net Configuration Hierarchy. Reprinted from [53]</i>	38
<i>Figure 13. 3D Cascade Architecture. Retrieved from [57].</i>	39
<i>Figure 14. Data Augmentation Example</i>	41
<i>Figure 15. Conceptual Framework.</i>	51
<i>Figure 16. CT scan in different views</i>	57
<i>Figure 17. nnU-Net V2 Architecture</i>	69

LIST OF TABLES

Table 1. Dataset Fingerprint Characteristics.....	61
Table 2. Data Augmentation Strategy.....	67
Table 3. Technical Specifications.....	73
Table 4. Training Hyperparameters	73
Table 5. Class Mapping	76
Table 6. Domain Expert (Radiologist) Profile.....	77
Table 7. Baseline Metrics for Kidney Abnormality Segmentation.....	83
Table 8. Metric Computation.....	89

Chapter 1

INTRODUCTION

This chapter establishes the foundation for a comparative investigation of single-class versus multi-class segmentation approaches for kidney abnormalities in CT imaging using nnU-Net V2. It begins by presenting the epidemiological context of renal pathologies and the critical role of CT imaging in diagnosis, then identifies the key computational gap: the lack of rigorous statistical comparison between single-class and multi-class deep learning segmentation models. The research questions and objectives are formulated to address this gap through a two-phase experimental design. In Phase 1, three single-class nnU-Net V2 models — each independently trained on one pathology (cyst, stone, or tumor) — are validated against published benchmarks to establish them as strong, competitive reference models. In Phase 2, a multi-class nnU-Net V2 model trained simultaneously on all three pathologies is compared directly against those validated single-class models to determine whether joint training improves, maintains, or degrades per-class segmentation accuracy.

1.1. Context of the Study

Kidney abnormalities represent a significant global health burden, encompassing a spectrum of pathologies ranging from benign cystic formations to life-threatening malignancies and painful stone formations [1] [2] [3]. Renal cysts are among the most frequent findings in adult imaging, with prevalence estimated at approximately 40% in adults over 50 years of age, following a distinct age-dependent increase from less than 1% in individuals aged 15-29 years to exceeding 22% in patients over 70 years [1]. Although the majority of these cysts are asymptomatic and benign, their high frequency necessitates

precise characterization to distinguish simple cysts from complex lesions that may harbor malignancy. Renal tumors represent a more critical diagnostic challenge, with greater than 50% of renal masses now detected incidentally during imaging performed for unrelated conditions, representing a dramatic shift from historical presentation patterns [4]. The incidence of renal cell carcinoma has been steadily increasing over the past four decades, partly driven by increased utilization of cross-sectional imaging that detects smaller, incidental tumors [4] [5]. Furthermore, kidney stones affect approximately 10% of the global population, with 106 million new cases reported worldwide in 2021 alone, reflecting a 27% rise since 2000 [3]. The presence of asymptomatic stone burdens can predict future symptomatic events, making early detection and quantification essential for preventative care and treatment planning [3] [6].

Computed tomography (CT) has established itself as the cornerstone imaging modality for evaluating kidney pathologies due to its superior spatial resolution, rapid acquisition times, and ability to characterize tissue properties through Hounsfield unit measurements [6]. For kidney stone detection, CT demonstrates exceptional diagnostic performance with sensitivity approaching 97% and specificity exceeding 95%, significantly outperforming traditional plain radiography and ultrasound [6]. The non-contrast CT protocol remains the gold standard for stone disease evaluation, as it allows for accurate stone detection, size measurement, and density assessment without the need for contrast administration [6]. In the context of renal masses, contrast-enhanced CT plays a pivotal role in differentiating benign cysts from solid tumors through multiphasic imaging protocols [7] [8]. The ability to assess enhancement patterns during arterial, venous, and delayed phases enables radiologists to characterize lesions according to the Bosniak

classification system for cystic masses and to evaluate the vascularity and extent of solid tumors [7]. Modern multidetector CT scanners provide isotropic voxel acquisition, facilitating high-quality multiplanar reconstructions and 3D visualization that aid in surgical planning and tumor staging. Despite these technological advantages, the interpretation of CT scans presents significant challenges. Radiologists must scrutinize hundreds of axial slices per examination, with each slice potentially containing subtle findings that require careful attention. The high volume of imaging studies in contemporary radiology practice, combined with the increasing complexity of cases, creates an environment prone to perceptual and cognitive errors, with burnout prevalence rates exceeding 54% among radiologists [9] [10].

Medical image segmentation is the computational process of partitioning medical images into distinct regions or objects of interest, enabling precise identification and delineation of anatomical structures and pathological abnormalities [11] [12]. It serves as a fundamental task in medical image analysis, transforming raw imaging data into structured, quantifiable information that supports clinical diagnosis, treatment planning, and disease monitoring [11]. In computed tomography imaging of the kidneys, segmentation enables radiologists to accurately measure lesion volumes, assess tumor boundaries for surgical planning, and quantify stone burden for treatment decisions [13]. The clinical importance of accurate kidney abnormality detection is underscored by documented diagnostic challenges. Studies have shown that patients frequently present with multiple concurrent kidney abnormalities requiring comprehensive evaluation [14]. These missed diagnoses or delayed characterizations can have profound clinical implications, as early detection of small renal cell carcinomas when still organ-confined is

associated with substantially better survival outcomes compared to advanced disease [5]. The increasing detection of small renal masses has intensified the clinical challenge of balancing intervention risks against cancer-related outcomes, emphasizing the need for precise imaging characterization to support evidence-based management decisions [4] [5].

The integration of artificial intelligence into medical imaging offers a potential solution to these diagnostic challenges. Deep learning frameworks, particularly Convolutional Neural Networks (CNNs), have demonstrated remarkable capability in automated analysis of medical images by learning hierarchical feature representations directly from data [11] [15]. In the domain of kidney abnormality detection, researchers have explored various computational approaches ranging from traditional machine learning to state-of-the-art deep learning architectures.

Previous work [16] developed an optimized YOLOv5 model for kidney stone detection in CT scans that incorporated Squeeze-and-Excitation blocks to improve feature extraction in noisy, low-dose images. Their approach achieved mean Average Precision (mAP) scores exceeding 90% on clean datasets, demonstrating the potential of real-time detection systems. Their model predicts stone locations based on high-density pixel clusters, attenuation values exceeding 400 Hounsfield Units, and spatial patterns within the renal collecting system. Through their model, they observed that detection accuracy remains high for stones greater than 5 mm but degrades significantly for smaller stones less than 3 mm. However, the bounding box outputs provide only coarse localization without precise lesion boundaries, limiting their utility for volumetric quantification and surgical planning [16].

In [17], extended the detection approach to segmentation, this time focusing on renal tumors. The authors implemented a 3D U-Net architecture for automated tumor segmentation on CT scans. Their model used volumetric features including tumor density, enhancement patterns, boundary characteristics, spatial location within the kidney, size measurements, and shape irregularity indices. Through the model, they revealed a critical performance disparity based on lesion size: achieving high accuracy for large exophytic masses but struggling with endophytic tumors embedded within the renal parenchyma, with tumors exhibiting heterogeneous enhancement patterns resulting in inconsistent segmentation boundaries. This size-dependent and morphology-dependent performance gap is particularly problematic because small and complex lesions are precisely those where accurate characterization offers the best clinical value [17].

The robustness of nnU-Net, a self-configuring segmentation framework, was also evaluated on external validation datasets from different hospital institutions to determine the generalizability of models trained on single-institution data when applied to multi-institutional cohorts [18]. Testing was conducted using an nnU-Net configuration based on the KiTS19 challenge framework [19]. Results show that while the model-maintained Dice scores around 0.85 on internal validation data from the training institution, segmentation accuracy degraded substantially when applied to CT scans from different scanner manufacturers and imaging protocols, with Dice scores dropping to 0.66 [18]. This cross-institutional performance drop was particularly pronounced for complex cystic lesions and small solid masses, indicating that domain shift—variations in scanner hardware, reconstruction kernels, and contrast timing—creates data distributions that differ from the training set, causing model failure in real-world deployment scenarios [18].

1.2. The Problem, Gap or Opportunity

The above-mentioned studies demonstrate significant progress in automated kidney abnormality detection but reveal critical gaps that limit clinical deployment and comprehensive diagnostic capability.

The first major limitation emerges from the predominant use of object detection methods that provide only coarse spatial localization. A recent evaluation [16] demonstrated a YOLO-based approach that achieves impressive detection performance with superior precision, recall, and mean Average Precision metrics compared to standard YOLOv5 variants. This approach demonstrates real-time processing capabilities suitable for clinical screening workflows. The integration of Squeeze-and-Excitation blocks specifically enhanced the model's ability to recalibrate channel-wise dependencies and capture intricate feature relationships, improving detection accuracy for small and irregularly shaped kidney stones [16]. However, the fundamental constraint of bounding box outputs means that these detection systems can only indicate the approximate region where a stone exists, without delineating its precise boundaries or three-dimensional morphology. This limitation becomes clinically significant when considering that accurate stone volume measurements are essential for treatment planning. Lithotripsy candidacy depends on precise stone size quantification, and longitudinal monitoring of stone burden requires volumetric rather than linear measurements [13]. Similarly, for renal masses, bounding boxes cannot distinguish between the actual tumor extent and surrounding normal parenchyma, limiting their utility for surgical planning where precise tumor margins determine the feasibility of nephron-sparing procedures versus radical nephrectomy [5]. The lack of pixel-level precision in detection-only approaches

fundamentally restricts their application to screening and triage functions, requiring subsequent manual segmentation by radiologists for any clinical decision that depends on accurate volumetric or morphological assessment.

The second critical gap involves morphology-dependent and anatomical complexity challenges in existing segmentation approaches. 3D U-Net combined with ResNet algorithms represents an advancement over detection-only methods by providing pixel-level boundary delineation for automatic renal mass segmentation and classification [17]. Their comprehensive study analyzed diverse tumor presentations and achieved promising overall performance metrics. However, segmentation of complex renal masses remains challenging due to substantial variability in tumor characteristics including size heterogeneity, location diversity (cortical versus medullary, upper pole versus lower pole), morphological complexity (exophytic protrusions versus endophytic embedding within parenchyma), and enhancement pattern variations that affect boundary clarity [17]. These morphology-dependent challenges are particularly problematic in clinical practice because the most diagnostically challenging cases are precisely the scenarios where accurate automated characterization would provide the greatest clinical value. Current segmentation systems that perform well on straightforward cases with clear boundaries may struggle with the complex presentations that cause highest inter-observer variability among radiologists and where accurate volumetric assessment is most critical for determining appropriate management strategies [5] [8].

The third fundamental limitation in current medical image segmentation is the pervasive reliance on single-institution datasets, which inherently restricts model generalizability. Deep learning architectures, while powerful, are highly susceptible to

domain shift—a phenomenon where a model overfits to the specific imaging characteristics of the hardware, reconstruction kernels, and contrast protocols of a single hospital. While some studies attempt to build local, private datasets, these single-center models frequently suffer catastrophic performance degradation when deployed in external environments.

Beyond these technical limitations, a fundamental research gap exists in the computer science literature: no study has conducted a rigorous, controlled comparison between single-class and multi-class segmentation models using the same architecture, dataset, and evaluation metrics. Existing studies that employ multi-class frameworks do not isolate or measure whether the joint training of multiple pathology classes improves, maintains, or degrades per-class segmentation performance compared to dedicated single-class models. This absence of controlled comparative evidence is significant — clinical deployment decisions about whether to use one unified model or separate specialized models currently lack empirical statistical grounding. Studies using nnU-Net V2 for kidney segmentation have demonstrated strong performance in tumor-focused single-class settings [18][19], and the MSWAL dataset has introduced multi-class ground truth annotations [20], yet no work has directly compared single-class versus multi-class nnU-Net V2 performance on the same kidney abnormality cases. This gap motivates the present study.

1.3. Research Questions

The aforementioned gaps and opportunities motivate the need to investigate comprehensive multi-class segmentation for kidney abnormalities. The research questions are formulated to address the following:

- (1) Do three independently trained single-class nnU-Net V2 models — one per abnormality (cyst, stone, tumor) — produce per-class segmentation performance

consistent with published state-of-the-art benchmarks [24][39][60][63][65], as measured by DSC, IoU, HD95, precision, recall, and F1-score, thereby establishing them as strong and validated reference models?

- (2) Does a multi-class nnU-Net V2 model trained simultaneously on cysts, stones, and tumors achieve statistically superior, equivalent, or inferior per-class segmentation performance compared to the three validated single-class nnU-Net V2 models?
- (3) What are the computational efficiency trade-offs including inference latency and GPU memory utilization between running three sequential single-class configurations versus one unified multi-class configuration, and do these trade-offs affect practical clinical deployment feasibility?

1.4. Research Objectives

To address the limitations of current kidney abnormality detection methods and extend studies in medical image segmentation, this study aims to develop and evaluate a multi-class segmentation system for kidney abnormalities in CT imaging using nnU-Net V2. This is achieved through a two-phase experimental design. Specifically, this study attempts to:

- (1) Train and evaluate three single-class nnU-Net V2 models — one independently optimized per abnormality class (cyst, stone, tumor) — and validate their per-class performance against published benchmarks [24][39][60][63][65] using DSC, IoU, HD95, precision, recall, and F1-score. This phase confirms that each single-class model constitutes a strong, competitive reference before any multi-class comparison is made.

- (2) Train and evaluate a multi-class nnU-Net V2 model on all three abnormality classes simultaneously and compare its per-class performance directly against the three validated single-class nnU-Net V2 models from Objective 1 to determine whether joint multi-class training improves, maintains, or degrades per-class segmentation accuracy. This is the primary comparison of the study.
- (3) Measure and compare computational efficiency — inference latency (seconds/volume) and peak GPU VRAM usage — between running the three single-class configurations sequentially versus the single multi-class configuration, to assess clinical deployment feasibility.

1.5 Significance of the Study

This study hopes to contribute to the advancement of automated medical image analysis for kidney abnormality detection, addressing critical gaps in diagnostic accuracy and clinical workflow efficiency. The significance of this research extends to multiple beneficiaries in both the medical and computer science communities.

For Radiologists and Clinical Practice. This study provides an automated tool that can assist radiologists in detecting kidney abnormalities with greater consistency and accuracy. Given the high volume of imaging studies and documented radiologist burnout rates exceeding 54% [9] [10], the developed system serves as a computational second reader that can flag potential abnormalities for radiologist review. The precise segmentation masks enable accurate volumetric quantification essential for treatment planning, tumor staging, and monitoring disease progression. By reducing diagnostic variability and improving detection of lesions across varying sizes and morphologies, this research directly contributes to earlier detection and improved patient outcomes [5].

For Computer Science and AI Research. This study advances the field of medical image segmentation by providing the first controlled statistical comparison between single-class and multi-class nnU-Net V2 configurations on kidney abnormality data. While single-class segmentation models have been extensively studied in isolation, and multi-class frameworks have been proposed for clinical comprehensiveness, no prior work has directly quantified the performance trade-off of these two paradigms under equivalent experimental conditions. The per-class comparison across cyst, stone, and tumor segmentation — with standardized metrics including DSC, IoU, HD95, precision, recall, and F1-score — provides empirically grounded evidence to guide model architecture decisions in medical AI deployment. Findings regarding whether multi-class training helps or hurts per-class performance may inform broader computer science research on multi-task learning trade-offs in medical segmentation tasks, contributing to the design of more efficient and accurate diagnostic AI systems.

For Healthcare Systems and Public Health. The automated segmentation system developed in this study has the potential to improve healthcare delivery efficiency by reducing the time required for image interpretation and standardizing abnormality detection across diverse imaging protocols and equipment [18]. By enabling earlier detection of kidney abnormalities and more accurate quantification, the system supports preventive care strategies and optimal treatment timing. The multi-class capability mirrors real clinical workflows where patients frequently present with multiple concurrent findings requiring comprehensive evaluation [14], making the system practical for deployment in hospitals and imaging centers.

1.6 Scope and Limitations (or Delimitations)

Since computed tomography is the primary imaging modality for comprehensive renal evaluation, offering superior spatial resolution and tissue characterization capabilities, this study focuses on developing a multi-class segmentation model for kidney abnormalities in CT imaging using nnU-NetV2. The model centers on the identification and delineation of four anatomical and pathological classes: kidneys (normal renal parenchyma), renal cysts (both simple and complex cystic lesions), renal tumors (solid masses including renal cell carcinoma), and kidney stones (nephrolithiasis). These abnormalities were selected based on their clinical prevalence—cysts affecting approximately 40% of adults over 50, tumors representing a growing incidental detection burden with over 50% now discovered asymptotically, and stones affecting 10% of the population [1] [3] [4]—and their critical importance for early detection and intervention.

This study executes four nnU-Net V2 training sessions across two phases. Phase 1 trains three single-class nnU-Net V2 models — one each for cysts, stones, and tumors — and validates each against published state-of-the-art benchmarks [24][39][60][63][65]. Phase 2 trains one multi-class nnU-Net V2 model on all three classes simultaneously and compares its per-class performance directly against the Phase 1 models. Per-class performance is measured using Dice Similarity Coefficient, Intersection over Union, Hausdorff Distance 95th percentile, precision, recall, and F1-score. Computational efficiency metrics — inference latency and GPU VRAM usage — are recorded to enable deployment feasibility comparison.

The study does not include other urinary tract abnormalities such as ureteral stones, bladder pathologies, or vascular anomalies, as these fall outside the scope of intra-renal

abnormalities. The research excludes rare kidney pathologies to maintain focus on the most clinically prevalent and diagnostically challenging abnormalities. The study does not incorporate other imaging modalities such as MRI, ultrasound, or nuclear medicine studies, as CT represents the primary modality for stone detection and comprehensive renal mass evaluation [6]. Furthermore, the research focuses on segmentation accuracy in a controlled research environment and does not include deployment as a real-time clinical decision support system, regulatory approval processes, or integration with hospital Picture Archiving and Communication Systems (PACS), as these implementation aspects require separate validation studies beyond the scope of this thesis.

The dataset is strictly limited to pathological cases exhibiting at least one of the three target abnormalities (cysts, stones, or tumors). Healthy control scans (negative cases) are excluded from the training distribution to maximize the model's exposure to pathological features within the computational constraints. Although the curated dataset comprises 60 volumetric cases per pathological entity, each case yields approximately 300 2D slices through the extraction of all three standard anatomical planes (axial, coronal, and sagittal), ensuring the model is exposed to a substantially larger effective training corpus than the case count alone suggests. The training and validation data for this study are strictly delimited to highly curated, multi-institutional public repositories (such as the KiTS datasets). Local institutional data collection was deliberately excluded from the scope to avoid the well-documented pitfalls of single-center bias, ensuring the model trains on the maximum possible variance in scanner hardware and imaging protocols.

Chapter 2

REVIEW OF RELATED LITERATURE & STUDIES

This chapter reviews the literature on automated kidney abnormality detection and segmentation in medical imaging. Section 1 examines the clinical significance and imaging characteristics of kidney cysts, stones, and tumors. Section 2 explores object detection methodologies in urological imaging, while Section 3 traces the evolution of U-Net architectures for medical image segmentation. Section 4 presents the nnU-Net framework, an automated self-configuring system achieving state-of-the-art performance across diverse segmentation tasks. Section 5 identifies critical research gaps, including the absence of unified multi-class segmentation systems, limited geographic diversity in training data, and insufficient integration with clinical classification frameworks. These sections provide the foundation for developing comprehensive AI-powered diagnostic tools capable of simultaneously detecting and characterizing multiple kidney abnormality types.

2.1 Clinical Significance of Kidney Abnormalities

2.1.1 Imaging Characteristics of Renal Abnormalities

In computed tomography (CT) imaging, the automated differentiation of renal structures relies entirely on pixel-level attenuation values, measured in Hounsfield Units (HU), alongside morphological features such as shape, texture, and boundary definition [6]. Healthy renal parenchyma typically exhibits a baseline attenuation that increases significantly and uniformly during contrast-enhanced phases [7], [8]. In stark contrast, renal cysts present as well-defined, smooth-walled, homogeneous masses with water-like attenuation [21]. Crucially for algorithmic differentiation, simple cysts do not exhibit

enhancement when intravenous contrast is administered [7]. Renal tumors present a more complex segmentation challenge [22]. They are typically solid masses with heterogeneous attenuation, exhibiting irregular enhancement patterns that vary depending on the contrast phase [7], [8]. This requires the neural network to learn complex spatial and textural gradients to accurately delineate their often-ill-defined borders from adjacent healthy tissue [17].

Kidney stones (renal calculi) introduce a distinct class imbalance and object scale challenge within the segmentation pipeline [23]. Unlike the soft tissue characteristics of cysts and tumors, stones are highly radiopaque, presenting with extreme attenuation values frequently exceeding 400 HU [16]. While their high pixel density provides a sharp contrast gradient against the renal collecting system, their significantly smaller volumetric footprint compared to tumors and cysts requires a highly sensitive feature extraction mechanism to prevent false negatives or boundary overlap errors [16], [24]. Ultimately, the distinct HU signatures, enhancement behaviors, and spatial distributions of these classes form the foundational quantitative data that the deep learning architecture must map.

2.1.2 Imaging Characterization Systems and Classification

In clinical practice, standardized scoring systems such as the Bosniak classification for complex renal cysts [25], [26] and the R.E.N.A.L. nephrometry score for solid masses are utilized to categorize abnormalities and guide surgical intervention. However, the efficacy of these systems relies heavily on manual visual assessment and the precise estimation of volumetric boundaries. Manual annotation and morphological grading by radiologists are not only highly time-consuming but also prone to significant inter-observer variability, contributing heavily to diagnostic bottlenecks and physician burnout [9], [10].

The inherently subjective nature of these manual evaluations limits the consistency of differential diagnoses across different institutions. Consequently, there is a critical need for automated, AI-driven segmentation frameworks capable of providing objective, pixel-perfect volumetric delineations to standardize clinical assessments and alleviate the manual workload.

The clinical imperative to automatically detect and characterize kidney abnormalities has motivated the development of deep learning approaches that combine two complementary operations: efficient detection of regions of interest and precise boundary delineation of abnormal structures. Detection and characterization tasks require fundamentally different technical approaches—detection must rapidly identify where abnormalities exist within large imaging volumes, while characterization demands pixel-level precision to measure lesion extent, volume, and relationship to critical anatomical structures. The foundation for precise pixel-level characterization lies in segmentation architectures that provide exact anatomical boundaries, enabling the volumetric measurements and feature extraction necessary for clinical decision support.

2.2 Object Detection Approaches in Urology

2.2.1 Object Detection in Medical Imaging: Concepts and Approaches

Object detection represents a fundamental computer vision task that combines two critical operations: localizing objects within an image and classifying them into predefined categories. In the context of medical imaging, this capability enables automated systems to identify and locate anatomical structures, pathological findings, and abnormalities within diagnostic images without requiring manual annotation by clinicians [11]. Deep learning approaches have revolutionized object detection by learning hierarchical feature

representations directly from data, enabling systems to identify subtle patterns that may indicate kidney abnormalities such as cysts, tumors, or structural anomalies [15]. Recent advances in imaging-based deep learning for kidney diseases have demonstrated broad applicability across clinical scenarios including organ segmentation, lesion detection, differential diagnosis, surgical planning, and prognosis prediction [14]. Contemporary applications extend beyond simple abnormality detection to encompass comprehensive analysis pipelines that support clinical decision-making through quantitative assessment and standardized reporting [13].

2.2.2 Detection vs. Segmentation: Complementary Methodologies

While object detection and segmentation are both essential components of medical image analysis, they serve distinct yet complementary purposes in the diagnostic pipeline. Object detection focuses on identifying the presence and location of abnormalities by generating bounding boxes that encompass regions of interest, providing spatial coordinates and class labels for detected objects [27]. In contrast, as illustrated in Figure 6, segmentation extends beyond simple localization to delineate precise pixel-level boundaries.




Figure 1. Comparison of Object Detection versus Semantic Segmentation in Medical Imaging. Adapted from [28].

These methodologies operate synergistically in urological imaging workflows, where detection algorithms first identify suspicious regions that may contain kidney abnormalities, subsequently guiding segmentation algorithms to perform detailed boundary delineation of the detected structures [11]. This complementary relationship proves particularly valuable in kidney abnormality analysis, where initial detection can rapidly screen entire imaging volumes to locate potential lesions, followed by precise segmentation that characterizes the abnormality's shape, size, and spatial relationship to

surrounding tissues. Recent studies have demonstrated that this hierarchical approach significantly improves computational efficiency while maintaining diagnostic accuracy [29].

The detection-then-segmentation paradigm enhances computational efficiency by focusing detailed segmentation analysis only on regions flagged as potentially abnormal, rather than processing entire image volumes at the highest resolution, thereby optimizing both diagnostic accuracy and processing time in clinical settings. This approach is particularly critical when dealing with three-dimensional volumetric data from CT or MRI, where exhaustive pixel-level processing would be computationally prohibitive [30].

2.2.3 Region-Based Deep Learning Detection Methods

Region-based convolutional neural networks have established a robust framework for object detection by decomposing the task into sequential stages of region proposal generation and object classification. The seminal Faster R-CNN architecture introduced an integrated Region Proposal Network that learns to generate high-quality candidate regions directly from image features, creating an end-to-end trainable detection system (illustrated in Figure 2) that has become foundational in medical imaging applications [15].




Figure 2. Faster R-CNN Architecture. Reprinted from [15].

Modern adaptations of region-based approaches incorporate attention mechanisms and feature pyramid networks to enhance multi-scale detection capabilities, enabling more effective identification of kidney abnormalities that vary substantially in size and appearance. These region-based approaches excel in medical imaging applications where detection accuracy takes precedence over processing speed, as the two-stage architecture allows for careful refinement of bounding box coordinates and classification decisions [31]. In urological imaging, region-based detectors demonstrate strength in identifying kidney abnormalities of varying sizes and morphologies, as the region proposal mechanism can adapt to detect both small focal lesions and larger masses within the same imaging study.

Recent implementations have integrated transformer-based feature extractors with region-based detection frameworks to capture long-range dependencies between

anatomical structures, improving the discrimination of pathological findings from normal anatomical variations [32]. However, their computational requirements remain higher compared to single-shot alternatives, making them more suitable for offline analysis rather than real-time clinical applications.

2.2.4 Single-Shot and Real-Time Detection Architectures

Single-shot detection architectures emerged to address the computational limitations of region-based methods by reformulating object detection as a unified regression problem that simultaneously predicts bounding box coordinates and class probabilities. The You Only Look Once (YOLO) family of detectors has evolved through multiple iterations, with recent versions incorporating advanced backbone networks and sophisticated feature aggregation strategies that maintain real-time processing capabilities while substantially improving detection accuracy for medical imaging applications (see Figure 3 for a typical YOLO architecture) [33].




Figure 3. YOLO Architecture. Reprinted from [34].

Contemporary YOLO applications in urological imaging have demonstrated remarkable success in kidney stone detection and kidney parenchyma localization. Recent

work presented an optimized YOLOv5 architecture specifically tailored for kidney stone detection in CT images, incorporating squeeze-and-excitation blocks within the C3 architecture to enhance channel-wise feature recalibration and achieve superior precision and recall metrics while maintaining real-time inference capabilities [16]. Similarly, YOLOv7 was applied to kidney detection in MRI, training on 878 patients with renal cell carcinoma and 206 with normal kidneys, demonstrating the architecture's adaptability to different imaging modalities and its capability to process medical image formats including DICOM and NIfTI [35]. While this study used MRI, the demonstrated capability to process both DICOM and NIfTI file formats validates YOLO's adaptability to diverse medical imaging modalities and preprocessing pipelines, directly relevant to CT-based workflows that similarly convert DICOM to NIfTI for volumetric processing.

Contemporary single-shot architectures employ feature pyramid networks and path aggregation mechanisms to enhance detection across multiple scales, improving the identification of kidney abnormalities with diverse sizes without sacrificing processing efficiency. Advanced focal loss mechanisms address class imbalance between foreground objects and background regions, a challenge particularly relevant in medical imaging where abnormalities constitute a small fraction of the total image area [36]. These single-shot architectures prove advantageous in clinical scenarios requiring rapid screening, such as intraoperative imaging or emergency radiology, where detection speed enables immediate clinical decision-making.

For urological applications, modern single-shot detectors facilitate efficient processing of multi-phase CT studies or large ultrasound video sequences, identifying kidney abnormalities across entire imaging series with minimal latency while achieving

detection performance increasingly comparable to region-based methods for most clinical tasks. The lightweight nature of these architectures makes them particularly suitable for deployment in resource-constrained clinical environments or edge computing scenarios where computational resources are limited [36].

2.2.5 Transformer and Advanced Detection Models

The emergence of transformer architectures has introduced novel attention mechanisms to object detection, enabling models to capture global contextual relationships that traditional convolutional approaches may overlook. Vision Transformers apply self-attention mechanisms to image patches, allowing networks to model long-range dependencies between spatially distant regions, which proves valuable when kidney abnormalities exhibit subtle imaging characteristics that require contextual interpretation [37].

Detection frameworks incorporating transformer architectures, inspired by models such as Detection Transformer (DETR), eliminate the need for hand-crafted components such as anchor boxes or non-maximum suppression by directly predicting object locations and classes through learned queries that attend to relevant image features (as shown in Figure 4) [38].




Figure 4. Vision Transformers. Reprinted from [38].

These transformer-based approaches demonstrate particular promise in medical imaging by leveraging attention mechanisms to focus computational resources on diagnostically relevant regions while maintaining awareness of broader anatomical context [31].

Recent applications in kidney imaging have explored dual-task approaches that combine detection with segmentation using transformer architectures. A transformer-based dual-task framework was presented for kidney MR segmentation in autosomal-dominant polycystic kidney disease, demonstrating the architecture's capacity to handle complex morphological variations while maintaining computational efficiency [30]. Transformers showed superior ability to capture global spatial relationships essential for distinguishing pathological from normal tissue in challenging cases where local appearance alone is insufficient.

The ability of transformers to integrate information across entire imaging volumes proves advantageous for detecting kidney abnormalities that may be characterized by subtle texture changes or relationships to surrounding structures rather than dramatic intensity differences. A comprehensive review of visual transformer applications in renal image analysis, emphasizing their effectiveness in capturing long-range dependencies critical for accurate delineation of complex renal structures and tumors that vary in location, shape, and appearance [37]. However, ongoing research continues to enhance their computational efficiency for clinical deployment, as transformer architectures typically require more computational resources than traditional CNN-based approaches.

2.2.6 Detection and Segmentation Integration in Urological Clinical Workflows

Detection algorithms serve as an efficient first-stage filter that rapidly scans urological imaging studies to identify regions containing potential abnormalities, generating bounding boxes that define areas requiring detailed analysis [39]. This hierarchical detection-to-segmentation pipeline optimizes computational efficiency by concentrating intensive segmentation processing only on clinically relevant regions rather than exhaustively analyzing entire imaging volumes [40]. Recent studies have demonstrated that automated longitudinal assessment systems employing this paradigm can effectively track stone burden progression and measure volumetric changes across serial CT imaging with minimal human intervention [13]. This standardization is particularly critical when developing models trained on multi-institutional datasets, as detection-based region proposal can help normalize variability in acquisition parameters and imaging protocols across different clinical sites [18].

2.3 Evolution of U-Net in Medical Image Segmentation

2.3.1 What is U-Net

U-Net is a specialized convolutional neural network architecture developed to address challenges in biomedical image processing where training data is typically limited and pixel-level localization is critical [12]. The architecture builds upon the foundation of Fully Convolutional Networks (FCNs) with significant modifications optimized for medical imaging applications. Its fundamental innovation lies in a symmetric encoder-decoder architecture, as illustrated in Fig. 5, that combines contextual information from a contracting path with precise spatial information from an expanding path through skip connections [12].




Figure 5. The U-Net Architecture. Illustration of the contracting and expanding paths with skip connections. Reprinted from [12].

The architecture consists of two main components defining its characteristic U-shaped structure. The contracting path captures context by progressively downsampling the input image, reducing spatial dimensions while increasing the number of feature channels to learn increasingly abstract representations [12]. Conversely, the symmetric expanding path enables precise localization by progressively upsampling feature maps. The critical innovation involves skip connections that directly link features from the contracting path to corresponding layers in the expanding path at the same resolution. These connections preserve fine-grained spatial details lost during downsampling, enabling the network to make precise predictions about pixel-level class assignments essential for accurate medical image segmentation [12].

Recent reviews have established U-Net-based architectures as the dominant paradigm for medical image segmentation across diverse anatomical structures and imaging modalities, with the original architecture serving as the foundation for hundreds of subsequent refinements and adaptations [28]. The enduring success of U-Net stems from its elegant balance between local precision and global context, achieved through the encoder-decoder structure with skip connections that has proven robust across varied segmentation tasks.

2.3.2 Architectural Components and Design Philosophy

The U-Net architecture eliminates fully connected layers and operates entirely through convolutional and pooling operations, making it computationally efficient and capable of processing images of arbitrary size [12]. Each layer in the network contains multi-channel feature maps, utilizing 3x3 kernels and ReLU activation functions throughout both the contracting and expanding paths. The symmetric design ensures that the number of feature channels in the expanding path mirrors those in the contracting path, allowing rich contextual information to propagate to higher resolution layers [12].

One of the fundamental design principles is the ability to function effectively with limited training data through aggressive data augmentation strategies [12]. Rather than requiring thousands of annotated samples like traditional deep learning approaches, U-Net can be trained effectively on very few images by applying elastic deformations and other transformation-based augmentations to available training data. This characteristic is particularly valuable in biomedical applications where annotated datasets are expensive and time-consuming to acquire. Furthermore, weighted loss functions are employed to

handle challenging segmentation problems, such as the separation of touching cells of the same class, by assigning higher weights to boundary regions in the loss calculation [12].

The architectural simplicity of U-Net belies its computational sophistication. By eschewing fully connected layers in favor of pure convolution operations, the network achieves several advantages: computational efficiency that enables processing of full-resolution medical images, translation invariance that allows consistent segmentation regardless of object position, and the ability to generate pixel-wise classifications that preserve spatial relationships essential for medical diagnosis [28]. These design choices have proven remarkably prescient, as subsequent architectures incorporating more complex components have often achieved only marginal improvements while sacrificing computational efficiency.

2.3.3 Clinical Applications and Modality Coverage

The architecture has demonstrated effectiveness across diverse medical imaging modalities and anatomical structures [28], [29]. Originally validated on the segmentation of neuronal structures in electron microscopy and cell tracking in phase contrast microscopy [12], the model has successfully extended to computed tomography (CT), magnetic resonance imaging (MRI), ultrasound, and X-ray imaging (see Figure 6 for examples across modalities) [28], [29].




Figure 6. Modality Versatility

This versatility stems from a modality-agnostic design that performs equally well on 2D image slices and volumetric 3D data after appropriate architectural modifications.

Clinical utility extends to numerous organ systems and pathological conditions, including brain tumor segmentation in multimodal MRI, cardiac structures in echocardiography, liver lesions in CT imaging, and lung structures in chest radiography [28], [29]. Recent comprehensive reviews of deep learning in medical image segmentation highlight U-Net's dominance across these applications, with the architecture serving as the baseline against which novel methods are evaluated [28]. This broad applicability reflects the robustness of the core architecture and its adaptability to task-specific modifications.

In kidney imaging specifically, U-Net has been applied to simultaneous segmentation of kidneys, tumors, and cysts in contrast-enhanced CT imaging [19], [41]. Its computational efficiency allows for the segmentation of high-resolution 3D volumes in seconds, making integration into diagnostic workflows clinically feasible [12]. This rapid processing capability is particularly important for interventional procedures where real-time or near-real-time guidance is required.

2.3.4 U-Net Variants and Architectural Innovations

The success of the baseline architecture has prompted numerous refinements to handle increasingly complex segmentation tasks [28], [29]. One prominent variant, U-Net++, incorporates nested and dense skip pathways to bridge the semantic gap between

feature maps extracted at different levels of the encoder and decoder (illustrated in Figure 7) [42].




Figure 7. U-Net++ Architecture. Reprinted from [42].

By incorporating dense convolutional blocks at each skip connection node, U-Net++ reduces information loss during the decoding process and has demonstrated improved performance across multiple medical imaging datasets [42].

Attention-based variants have emerged as a major refinement direction, incorporating mechanisms to selectively focus on clinically relevant regions while suppressing background noise. Attention U-Net integrates attention gates within the skip connections to recalibrate feature maps based on spatial and channel-wise importance, improving segmentation accuracy particularly for small lesions (see the attention gate mechanism in Figure 8) [43].




Figure 8. Attention Mechanism. Reprinted from [43].

Recent implementations specifically for kidney tumor segmentation have demonstrated that boundary attention mechanisms can significantly improve segmentation accuracy at tumor edges, addressing one of the primary challenges in renal mass characterization [41]. MSS U-Net introduced multi-scale supervised learning for 3D kidney and tumor segmentation, training the network to produce accurate predictions at multiple resolution levels simultaneously [44]. Similarly, recent work on the KiTS23 challenge explored various 3D U-Net training configurations to achieve superior segmentation performance [45]. Furthermore, 3D U-Net extends the original architecture to volumetric medical imaging by replacing 2D convolutions with 3D convolutional operations [28], [29]. Recent studies employing 3D U-Net for renal mass segmentation have achieved clinical-grade performance [17]. Hybrid variants like ResUNet combine the encoder-decoder structure with residual learning principles [28]. TransUNet pioneered the transformer approach by using a Vision Transformer as the encoder while maintaining U-Net's decoder structure (as depicted in Figure 9) [46].




Figure 9. Hybrid Transformers. Reprinted from [46].

Swin-UNet further refined this concept using Swin Transformers with shifted window attention, demonstrating superior performance on multi-organ segmentation tasks including kidney structures [47]. The STC-UNet architecture incorporates Selective Kernel modules to address challenges posed by unclear boundaries and positional uncertainties in renal tumors [48]. Similarly, dual attention approaches have shown promise for ultrasound-based renal tumor segmentation, achieving MIOU scores exceeding 93% on clinical datasets [49].

2.3.5 Performance Metrics and Clinical Benchmarking

Medical image segmentation performance is evaluated using several complementary metrics that capture different aspects of segmentation quality [28]. The Dice Similarity Coefficient (DSC) measures the overlap between predicted and ground truth segmentation masks, providing a balanced measure that accounts for both false positives and false negatives. Intersection over Union (IoU), also known as the Jaccard index, computes the ratio of the intersection to the union of predicted and ground truth

regions, offering a stricter assessment of segmentation accuracy (see the visual comparison of DSC and IoU in Figure 10).




Figure 10. Volumetric Overlap Metrics

While these overlap-based metrics provide quantitative measures of volumetric accuracy, boundary precision is typically quantified using the Hausdorff distance and Average Symmetric Surface Distance [28].

Performance reporting typically includes sensitivity (recall), specificity, and precision to provide a comprehensive assessment of segmentation quality across different clinical requirements [28]. Sensitivity measures the proportion of actual positive pixels correctly identified, critical for ensuring complete lesion capture. Specificity quantifies the proportion of actual negative pixels correctly classified, important for minimizing false positive regions. Precision, or positive predictive value, indicates the proportion of predicted positive pixels that are actually positive, relevant for clinical confidence in segmented regions.

For clinical translation, published benchmarks suggest that Dice coefficients exceeding 0.85 for organ segmentation and 0.75 for lesion segmentation generally indicate

clinically acceptable performance [28], [29]. However, mathematical accuracy metrics do not always correlate perfectly with clinical utility, necessitating a careful balance between metric optimization and the improvement of patient outcomes [28]. Recent studies on kidney and tumor segmentation have reported achieving these benchmarks, with state-of-the-art models attaining Dice scores of 0.948 for kidney segmentation, 0.776 for masses, and 0.738 for tumors on standardized challenge datasets [45].

Surface-based metrics have gained prominence for assessing boundary accuracy, particularly important in surgical planning where precise delineation of tumor margins is critical. The 95th percentile Hausdorff Distance (HD95) measures the maximum distance between predicted and ground truth boundaries while being robust to outliers, with lower values indicating better boundary alignment (illustrated in Figure 11).




Figure 11. Boundary Distance Metrics. Reprinted from [50].

Recent kidney segmentation challenges have incorporated Surface Dice metrics that combine volumetric overlap with surface distance considerations, providing a more comprehensive assessment of segmentation quality for clinical applications [45].

The complementary relationship between detection and segmentation defines optimal clinical workflows for kidney abnormality analysis—detection serves as an efficient initial screening mechanism that identifies regions requiring detailed analysis, while segmentation provides the precise anatomical boundaries necessary for clinical measurement and characterization. Implementing this integrated pipeline in clinical practice, however, faces substantial challenges: manual configuration of detection and segmentation parameters is time-intensive and requires deep machine learning expertise; performance is highly dependent on dataset-specific parameter tuning and validation; cross-institutional robustness demands careful attention to diverse imaging protocols,

scanner manufacturers, and image acquisition parameters; and reproducibility across clinical sites is compromised by the sensitivity of both detection and segmentation networks to variations in training and configuration.

The nnU-Net framework addresses these implementation challenges through fully automated configuration, eliminating the need for manual task-specific parameter tuning while achieving state-of-the-art performance across diverse segmentation tasks. By automatically configuring U-Net architectures for kidney-specific tasks and encoding domain knowledge about optimal detection-segmentation integration into systematic rules, nnU-Net transforms kidney abnormality segmentation from a specialized research capability requiring substantial machine learning expertise to a clinically deployable tool accessible to clinical researchers and radiologists without formal deep learning training. This transformation represents a critical advance toward clinical translation of automated kidney abnormality analysis. The following section examines the nnU-Net framework's architecture, configuration strategy, empirical validation, and implications for clinical deployment of kidney abnormality detection and characterization systems.

2.4 The nnU-Net Framework

2.4.1 Overview of nnU-Net

nnU-Net is an automated self-configuring framework for deep learning-based medical image segmentation that has fundamentally transformed the clinical deployment of AI-powered diagnostic tools [51]. It automatically configures every component of the segmentation pipeline without requiring manual intervention beyond providing raw imaging data and corresponding binary segmentation annotations [51]. The fundamental motivation stems from the observation that small errors in configuration choices can lead

to large drops in performance [51]. Analysis revealed that method configuration quality often exceeds architectural novelty, motivating a systematic, automated process [51]. The framework achieved first place in the KiTS19 kidney tumor segmentation challenge, establishing a benchmark for multi-institutional CT kidney segmentation [19]. Subsequent iterations have maintained dominance while preserving the core automated configuration principles [51]. Recent comprehensive benchmarking against contemporary architectures including Vision Transformers and Mamba-based models has validated that properly configured CNN-based U-Net variants within the nnU-Net framework continue to achieve state-of-the-art performance when rigorously evaluated [52].

2.4.2 Automated Configuration Architecture

The framework implements automated configuration through a three-tier parameter hierarchy (shown in Figure 12) that systematically handles design decisions [51].




Figure 12. *nnU-Net Configuration Hierarchy. Reprinted from [53]*

The first tier consists of fixed parameters, including specific activation functions and loss function strategies combining Dice and cross-entropy loss [51]. The second tier comprises rule-based parameters, which establish explicit dependencies between the dataset fingerprint and the pipeline fingerprint [51]. One critical rule governs the interdependent configuration of batch size and patch size while respecting GPU memory constraints [51], [54]. Additional rules determine intensity normalization approaches

tailored to specific modalities such as CT Hounsfield unit clipping or MRI z-score normalization [51]. The third tier consists of empirical parameters covering decisions that benefit from data-driven optimization on the specific dataset [51]. This includes determination of post-processing strategies such as connected component analysis [51], [54]. This strategy avoids the computational overhead of exhaustive hyperparameter search [51], [55].

2.4.3 Cascade Architecture and Multi-Resolution Processing

A critical feature is the systematic use of cascaded U-Net configurations for anatomical structures that occupy relatively small regions within imaging volumes [51]. The framework automatically detects spatial thresholds, triggering a two-stage cascade approach [51], [56]. In this configuration, the first U-Net operates on downsampled images to perform coarse localization of the target structures, while the second U-Net operates at full resolution within regions of interest identified by the first stage (see the cascade workflow in Figure 13) [52].




Figure 13. 3D Cascade Architecture. Retrieved from [57].

This cascade trigger directly addresses the fundamental trade-off between capturing sufficient contextual information for accurate localization and maintaining sufficient

spatial resolution for precise boundary delineation [56], [58]. By automatically selecting cascade configurations, the system encodes domain knowledge about volumetric imaging without requiring manual intervention [51]. This design proves particularly valuable for kidney imaging applications where structures benefit substantially from the coarse-to-fine processing strategy [56], [58]. Recent implementations demonstrated the effectiveness of this approach for kidney tumor segmentation [56]. A two-stage cascaded framework effectively refined segmentation of kidney tumors and cysts [56]. This approach demonstrated that the cascade architecture effectively handles the multi-scale nature of kidney abnormality segmentation [56]. The intermediate cropping process maintains full-resolution detail for precise boundary characterization [56], [58].

2.4.4 Cross-Institutional Robustness and Domain Generalization

Robustness to domain shift represents a critical consideration for clinical deployment of automated segmentation systems [55]. The nnU-Net framework addresses domain shift through comprehensive data augmentation strategies scaled based on dataset properties [55]. Applied augmentations include geometric transformations and specialized techniques (an example of these augmentations are provided in Figure 14) [51], [55].




Figure 14. Data Augmentation Example

Empirical validation of cross-institutional robustness has been conducted across multiple imaging modalities and anatomical targets. nnU-Net's generalization capability has been demonstrated, showing that automated data augmentation maintained high Dice coefficients across different vendor domains comparable to performance within same-vendor data [55]. A rigorous evaluation of nnU-NetV2 was tested on an independent institutional dataset comprising 139 cases [18]. The study revealed models trained on diverse multi-institutional data demonstrated improved segmentation robustness and generalizability [18]. This finding emphasizes training dataset diversity maintains performance across diverse real-world imaging environments [18]. Models trained exclusively on single-institution data may fail to generalize to external populations, limiting practical utility in multi-site clinical networks [18]. The framework provides a pathway to supporting the deployment of AI-assisted diagnostic tools in real-world healthcare settings [18].

2.4.5 Empirical Performance and Validation

The empirical validation of nnU-Net spans an unprecedented diversity of segmentation tasks across modalities, anatomical structures, and imaging protocols [51]. Comprehensive evaluations have established nnU-Net as a state-of-the-art framework on the majority of tested public benchmarks [51]. The framework secured first place in multiple international biomedical imaging challenges [19], [51]. The KiTS19 challenge results showed robust performance across diverse tumor sizes and locations [19]. Subsequent challenges have validated nnU-Net's continued dominance, achieving high Dice scores for tumors on standardized challenge datasets [45], [59]. Comprehensive benchmarking against contemporary architectural alternatives including Vision Transformers and Mamba-based models has validated that properly configured CNN-based U-Net variants within the nnU-Net framework continue to achieve state-of-the-art performance when rigorously evaluated [52]. The study demonstrated that CNN-based approaches outperform alternatives on standard medical segmentation datasets [52]. Many reported performance improvements in recent literature resulted from validation shortcomings rather than genuine architectural superiority [52]. This validation confirms the approach remains highly competitive against novel architectures designed for large-scale datasets [52]. For kidney abnormality segmentation, these findings support the selection of nnU-Net over more computationally demanding Transformer alternatives [52], [60].

2.4.6 Clinical Translation and Accessibility

The accessibility of nnU-Net for clinical deployment represents a significant advance for hospital settings [51]. The framework requires minimal specialized

knowledge, allowing clinical teams to apply the method directly to raw medical imaging data and binary annotation masks [51], [61]. Standard institutional GPU resources are generally sufficient, avoiding the need for high-performance computing infrastructure or multi-GPU clusters [51], [61]. The deployment workflow executes automated training scripts [51], [53]. The framework automatically handles all configuration decisions through the three-tier parameter hierarchy [51]. Resulting models generate pixel-level segmentation maps suitable for treatment response assessment, and surgical planning [13], [51]. This accessibility transforms automated segmentation into a tool deployable without formal deep learning training [51]. Recent applications demonstrate successful deployment achieving clinically acceptable segmentation performance with reduced computational requirements [62]. Such optimizations expand the practical applicability of automated segmentation tools to a broader range of healthcare facilities [62]. The framework's ability to handle multi-class segmentation tasks positions it as the natural choice for comprehensive characterization encompassing tumors, cysts, and stones [45], [56]. The cross-institutional robustness supports deployment across diverse clinical imaging environments [18], [58].

2.4.7 State-of-the-Art Deep Learning Frameworks for Kidney Segmentation

Building upon the clinical accessibility of these models, recent comparative studies have firmly established the superiority of self-configuring frameworks over traditional architectures for renal image analysis. A controlled benchmarking study directly compared the nnU-Net framework against a custom 3D U-Net for kidney and tumor segmentation [60]. By maintaining identical preprocessing, data augmentation, and training protocols, it was demonstrated that nnU-Net provides a highly robust baseline, outperforming the

standard 3D U-Net in both Dice Similarity Coefficient (DSC) and Intersection over Union (IoU) [60]. Building upon these architectural advantages, the specific capabilities of the nnU-Net V2 architecture were investigated using a 3D full-resolution configuration [63]. The application of the model to the KiTS23 dataset yielded high DSC and IoU scores, cementing nnU-Net V2 as a highly optimized foundational architecture for multi-class renal segmentation tasks [63].

2.5 Identified Research Gaps

The comprehensive review of literature on kidney abnormality detection and segmentation reveals substantial progress in automated analysis of individual abnormality types, yet identifies critical gaps that limit the clinical applicability and comprehensiveness of existing approaches. These gaps motivate the development of integrated multi-class segmentation systems capable of simultaneously detecting and characterizing all major kidney abnormalities within unified clinical workflows.

2.5.1 Single-Class Focus in Existing Segmentation Systems

Current state-of-the-art deep learning approaches for kidney abnormality segmentation predominantly focus on single abnormality types, with separate specialized models developed for kidney tumors, cysts, or stones [18], [19], [39]. While individual models achieve impressive performance within their specific domains—tumor segmentation models reaching Dice scores exceeding 0.85 on KiTS challenge datasets [18], stone detection models achieving accuracies above 96% [64], and cyst segmentation showing robust performance in polycystic kidney disease applications [30]—no comprehensive system simultaneously addresses all three major abnormality categories within a unified framework.

This single-class paradigm creates substantial limitations for clinical deployment. In real-world diagnostic workflows, patients frequently present with multiple concurrent kidney abnormalities requiring comprehensive evaluation [14]. A patient may harbor both renal masses and incidental cysts, or demonstrate stone disease alongside cystic lesions. Current approaches necessitate running multiple independent models sequentially, each with its own preprocessing pipeline, inference requirements, and post-processing steps, resulting in computational inefficiency and workflow complexity that impedes clinical integration [45]. Furthermore, independent single-class models fail to leverage potential synergies in learning shared anatomical context and imaging features across abnormality types, potentially limiting both efficiency and performance.

The absence of multi-class kidney abnormality segmentation systems represents a fundamental gap in translating automated analysis from research benchmarks to comprehensive clinical decision support tools. Developing unified frameworks capable of simultaneous detection and precise delineation of tumors, cysts, and stones would substantially advance clinical applicability by providing radiologists with complete abnormality characterization from single imaging studies, supporting comprehensive diagnostic assessment and treatment planning [14].

2.5.2 Fragmented Public Datasets and the Need for Data Unification

While substantial public datasets exist for kidney abnormality segmentation, these resources remain fragmented across different clinical objectives and pathological entities. Datasets such as KiTS19, KiTS21, and KiTS23 focus primarily on tumor segmentation [19], [45], while other publicly available resources address specific pathologies like kidney stones or cystic lesions in isolation. This fragmentation creates a fundamental challenge:

existing models are trained to recognize and segment individual abnormality types, but no unified public dataset exists that enables comprehensive multi-class segmentation encompassing tumors, cysts, and stones simultaneously.

This data fragmentation has significant clinical implications. In real-world diagnostic radiology, patients frequently present with multiple concurrent kidney pathologies, a tumor-bearing kidney may also contain cysts or stones, yet current segmentation approaches trained on isolated datasets cannot comprehensively characterize such cases. Deploying separate models for each abnormality type introduces computational inefficiency, potential inconsistencies in overlapping or adjacent structures, and workflow complexity that limits clinical adoption. The absence of a unified training framework that exposes models to the full spectrum of kidney abnormalities during learning prevents development of truly comprehensive diagnostic systems.

The research opportunity lies in unifying these disparate public datasets into a single multi-class training framework. By combining tumor-focused datasets (KiTS series) with publicly available stone and cyst datasets, we can create a comprehensive training corpus that exposes the model to the full range of kidney pathologies simultaneously. This unified approach addresses a fundamental gap in the current research landscape: while individual public datasets exist for specific pathologies, no prior work has systematically integrated these resources to enable simultaneous detection and segmentation of all major kidney abnormality types within a single model architecture.

This data unification strategy offers several advantages beyond simply training separate models on each dataset. Multi-class training enables the model to learn shared anatomical features and spatial relationships between different pathology types, potentially

improving overall segmentation accuracy through cross-pathology feature learning. Additionally, a unified model reduces computational overhead during inference and simplifies clinical deployment compared to maintaining multiple specialized systems. Our study bridges this critical gap by developing the first comprehensive multi-class kidney abnormality segmentation framework trained on unified public datasets, demonstrating that integration of disparate data sources can produce clinically viable comprehensive diagnostic systems [18], [45].

Recent advancements have begun to address this data fragmentation. The introduction of the 3D Multi-class Segmentation of Whole Abdominal Lesions (MSWAL) dataset represents a significant shift toward unified pathological annotations, providing comprehensive voxel-level ground truth for multiple abdominal abnormalities, including kidney stones, cysts, and tumors, within single volumetric scans [20]. While MSWAL provides a robust foundation for multi-class learning, it is constrained by its single-institution origin. Therefore, this study's strategy of harmonizing a multi-class dataset (MSWAL) with a multi-institutional dataset (KiTS23) address both the structural fragmentation of public datasets and the critical need for cross-institutional robustness in model training.

2.5.3 Insufficient Integration of Clinical Classification Systems

While automated segmentation provides precise anatomical boundaries of kidney abnormalities, existing systems generally do not integrate outputs with established clinical classification frameworks that guide diagnostic and therapeutic decisions. The Bosniak classification for cystic lesions, R.E.N.A.L. nephrometry scoring, and stone burden quantification metrics remain underdeveloped in automated computation [5], [7]. Accurate

segmentation alone is insufficient without Bosniak category assignment based on wall thickness, septation characteristics, enhancement patterns, and calcification presence [7]. Renal mass segmentation requires exophytic properties and polar location for surgical planning and intervention selection [5]. Stone disease management depends critically on quantitative metrics including total stone burden, individual dimensions, and density characteristics that inform composition and treatment susceptibility [16].

2.5.4 The Comprehensive Metric Evaluation Gap

In addition to the limitations in dataset fragmentation and clinical integration, a review of the current baseline literature reveals a significant gap in comprehensive, multi-class clinical evaluations. Beyond standard overlap metrics like DSC and IoU, rigorous clinical validation requires assessing boundary distances and prediction exactness. A comprehensive validation study of an AI-based kidney segmentation framework was recently conducted utilizing evaluation metrics such as the 95th percentile Hausdorff Distance (HD95), Precision, Recall, and F1-score [65]. While this provides a strong methodological baseline for multi-metric evaluation, a notable limitation is the consolidation of distinct pathological classes. In this framework, cysts and tumors were merged into a singular "abnormality" class, inherently limiting the model's utility for precise differential diagnosis [65].

Furthermore, while foundational studies establish standard performance benchmarks for kidneys and tumors [60], [63], none of the current state-of-the-art baselines incorporate kidney stones alongside cysts and tumors in their evaluation. Consequently, there remains an unmet need for a single, unified framework capable of simultaneously isolating distinct renal pathologies—specifically kidneys, tumors, cysts, and stones—while

being evaluated against a comprehensive suite of overlap, boundary, and efficiency metrics. To address the absence of kidney stones in unified multi-class evaluations, baseline metrics for this specific pathology must be drawn from dedicated stone segmentation models. Recent automated systems have successfully established volumetric segmentation baselines for kidney stones on non-contrast CT scans [24]. Furthermore, specialized deep segmentation networks have been utilized to simultaneously segment both the renal structure and calculi, providing essential overlap metrics such as DSC for stone detection [39]. In this study, the performance of the proposed multi-class nnU-Net framework on the stone class will be benchmarked against these dedicated single-task and dual-task models to ensure rigorous clinical validation across all isolated abnormalities.

2.5.5 Gap Summary and Research Opportunity

The identified gaps collectively point to a specific, addressable research opportunity in computer science: the absence of a controlled comparative study between single-class and multi-class nnU-Net V2 segmentation for kidney abnormalities. While prior work has validated nnU-Net V2 for kidney tumor segmentation [18][19] and the MSWAL dataset enables multi-class ground truth annotation [20], no study has isolated the performance effect of training strategy — single-class versus multi-class — using the same architecture, data, and metrics. This study addresses that gap directly, using 3D U-Net as a consistent baseline reference, and provides the first statistically grounded comparison to guide practitioners in choosing between single-class and multi-class segmentation pipelines for renal CT analysis.

Chapter 3

METHODOLOGY

This chapter presents the comprehensive methodological framework and computational procedures developed to achieve automated, high-precision 3D segmentation of renal structures and pathologies. Grounded in a retrospective analysis of clinical Computed Tomography (CT) data, the study utilizes the self-configuring nnU-Net v2 deep learning architecture to identify and delineate kidneys, cysts, stones, and tumors. The methodology is organized through an Input-Process-Output (IPO) model, spanning from physiological normalization of raw radiological data to the rigorous quantitative evaluation and qualitative clinical review by radiologists.

3.1 Conceptual Framework

The methodological architecture of this research is grounded in the intersection of computational medical imaging, deep learning-based semantic segmentation, and clinical decision support system. The primary objective is to translate raw radiological data into actionable clinical insights through the automated delineation of kidney substructures and pathologies – specifically cysts, stones, and tumors. This study adopts a rigorous Input-Process-Output (IPO) model, a fundamental system engineering approach that structures the methodology into three distinct but interconnected phases: data ingestion and preparation (Input), computational analysis and model training (Process), and performance evaluation and clinical inference (Output).

3.1.1 The Input-Process-Output (IPO) Model

The IPO model serves as the architectural skeleton of this thesis, ensuring a logical flow of data and rigorous validation of the segmentation results (as illustrated in the framework diagram in Figure 20).




Figure 15. Conceptual Framework

This system-theoretic perspective allows for the isolation of variable at each stage, ensuring that errors can be traced back to specific components – whether they

be artifacts in the input data, convergence failures in the process, or metric misinterpretations in the output.

3.1.2 Input Phase: Data Ingestion and Physiological Standardization

The foundation of the framework lies in the acquisition of high-fidelity volumetric CT data. The input stage is characterized by the transformation of raw medical data into a machine-readable format. This involves the ingestion of DICOM (Digital Imaging and Communications in Medicine) series, which contain not only the pixel data but also critical metadata regarding the acquisition parameters. The conversion of these series into the NIfTI (Neuroimaging Informatics Technology Initiative) standard is a pivotal step, as it creates a unified volumetric representation that preserves the affine matrix—defining the spatial relationship between the image voxels and the physical world coordinate system.

The preprocessing phase within the Input stage is conceptually critical as it bridges the domain gap between physical tissue density—measured in Hounsfield Units (HU)—and the numerical inputs required by the neural network. By applying specific intensity clipping (Focus range of -135 to 215 HU) and resampling techniques (1.0 mm isotropic spacing), the input stage standardizes the physical variance across different patients and scanners. This "physiological normalization" ensures that the neural network learns to associate specific numerical ranges with specific tissue types (e.g., distinguishing a dense stone from soft tissue), rather than learning scanner-specific noise patterns.

3.1.3 Process Phase: Deep Learning and Volumetric Segmentation

The process component is driven by the 3D U-Net architecture within the nnU-Net v2 framework. The conceptual novelty here is the use of deep supervision, where the network learns to segment structures at multiple resolutions simultaneously. This mechanism involves injecting auxiliary loss layers at various stages of the decoder pathway, thereby improving gradient flow deep into the network and maximizing feature extraction from both coarse and fine spatial scales.

This stage encompasses the iterative training of the model using a composite loss function (Dice Loss plus Cross-Entropy Loss), stochastic optimization via Stochastic Gradient Descent (SGD) with Nesterov momentum, and massive on-the-fly data augmentation. The augmentation strategy is designed to enforce translational, rotational, and intensity invariance, forcing the model to learn robust anatomical features rather than memorizing the specific orientation or contrast levels of the training cases. The use of 5-fold cross-validation during this phase ensures that the model's performance is robust across the entire dataset distribution and not biased by a specific train-test split.

3.1.4 Output Phase: Evaluation and Clinical Inference

The output of the framework is a multi-class segmentation map where each voxel is classified into one of five categories: background, kidney, cyst, stone, or tumor. The conceptual validity of these outputs is assessed using a dual-metric approach:

1. **Volumetric Overlap:** The Dice Similarity Coefficient (DSC) quantifies the spatial overlap between the prediction and the ground truth, providing a global measure of segmentation quality.
2. **Boundary Accuracy:** The Hausdorff Distance 95th Percentile (HD95) measures the worst-case error at the segmentation boundaries, offering critical insight into the geometric precision of the model, which is vital for surgical planning and radiotherapy.

Furthermore, the framework incorporates a "Human-in-the-Loop" validation mechanism, where expert radiologists review failure cases (identified by low DSC scores) to provide qualitative feedback. This step transforms the thesis from a purely computational exercise into a clinically relevant study, categorizing errors into detection failures, false positives, or boundary inaccuracies.

3.2 Data Management

The integrity, diversity, and quality of the dataset are paramount to the success of any deep learning model, particularly in the medical domain where data scarcity and class imbalance are prevalent challenges. The data acquisition strategy for this thesis is designed to ensure strict compliance with medical data interoperability standards, ethical guidelines regarding patient privacy, and the specific technical requirements of the nnU-Net v2 training pipeline.

3.2.1 Sources and Ethical Considerations

The dataset was curated from publicly accessible biomedical imaging repositories and educational resources. All renal abnormality cases were obtained from established online medical imaging databases and publicly available radiologic archives. Additional cases were supplemented using anonymized examples from open-access academic sources to ensure representative coverage of common renal pathologies. Given the retrospective nature of this study—involving the computational analysis of previously collected clinical data generated during routine patient care—strict ethical protocols are observed to protect subject confidentiality. Each case represents a complete 3D volumetric CT scan comprising approximately 100 slices per anatomical plane. By extracting all three standard views — axial, coronal, and sagittal — each volumetric case yields approximately 300 2D slices, resulting in approximately 18,000 individual 2D slices per pathological entity and 54,000 slices across the full dataset.

Data Selection Strategy: To ensure the model generalizes across heterogeneous clinical presentations, a two-stage data curation protocol is employed. Initially, a candidate pool of 150 volumetric cases per pathological entity (Cysts, Stones, Tumors) will be aggregated from the sources. From this pool, a stratified selection process filters the data down to 60 high-value volumetric cases per entity (Total of 180 unique pathological volumes). Cases are selected based on a Diversity Criteria to maximize variance in:

1. Size: Balancing micro-lesions ($<5\text{mm}$) with large masses ($>4\text{cm}$).

2. Location: Ensuring representation of upper, mid, and lower pole lesions, as well as cortical vs. medullary placement.
3. Morphology: Selecting cases with varying densities (HU values) and shapes (e.g., simple vs. complex cysts, staghorn vs. focal stones).

To ensure absolute methodological transparency, the experimental corpus is strictly defined. While the initial curation pipeline aggregated a candidate pool of 150 volumetric cases per pathological entity, the rigid inclusion and exclusion criteria filtered this pool down to a finalized experimental dataset of exactly 180 volumetric CT scans ($N = 180$), equally balanced across three pathological classes: 60 volumes featuring renal cysts, 60 featuring nephrolithiasis, and 60 featuring solid renal neoplasms. It is exclusively this 180-volume cohort that is subjected to the computational pipeline. During the automated 5-fold cross-validation, this corpus is partitioned such that each fold utilizes 144 volumes (80%) for training and 36 volumes (20%) for internal validation.

This curation ensures that although the final dataset size is compact, it is mathematically representative of the disease spectrum, preventing the model from over-optimizing on "easy" or repetitive cases.




Figure 16. CT scan in different views

3.2.2 Ethical Compliance Framework:

1. Retrospective Study Protocol: Since this study utilizes publicly available, pre-anonymized datasets distributed under CC-BY-NC 4.0, institutional IRB approval was not required. The study strictly adheres to the data usage agreements of the respective repositories.
2. De-identification and Anonymization: All patient data is strictly anonymized at the source prior to transfer for computational analysis. This process involves the irreversible removal or masking of all Protected Health Information (PHI) from the DICOM headers. Specific tags such as Patient Name (0010,0010), Patient ID (0010,0020), Patient Birth Date (0010,0030),

and Institution Name (0008,0080) are scrubbed or replaced with study-specific pseudonyms (e.g., Kidney_001). This ensures that the dataset used for training contains no linkages to the original patient identities, complying with data privacy regulations such as HIPAA or GDPR.

3. **Data Security:** The anonymized data is stored on secure, encrypted servers (e.g., Google Drive mounted to Colab Pro+ with restricted access), ensuring that only authorized research personnel have access to the imaging volumes.

3.2.3 DICOM to NIfTI Conversion

Raw medical images are generated by CT scanners in the DICOM (Digital Imaging and Communications in Medicine) format. While DICOM is the standard for clinical storage (PACS), it is suboptimal for deep learning due to its slice-based file structure (one file per slice) and complex metadata hierarchy. Deep learning frameworks, including nnU-Net, require the NIfTI (Neuroimaging Informatics Technology Initiative) format (.nii.gz), which encapsulates the entire 3D volume into a single compressed file and strictly defines the spatial coordinate system (affine matrix) necessary for 3D volumetric operations.

Conversion Protocol:

1. **Input:** A directory containing the series of DICOM slices for a single patient study.
2. **Processing:** dcm2niix reads the DICOM header information to reconstruct the contiguous 3D volume. It validates the slice ordering based on the

ImagePositionPatient tag and handles any gantry tilt correction to ensure the resulting volume is orthogonal.

3. **Metadata Extraction:** In addition to the image volume, dcm2niix generates a JSON sidecar file containing critical acquisition parameters such as slice thickness, pixel spacing, manufacturer, and convolution kernel. This metadata is essential for the subsequent "Dataset Fingerprinting" stage of nnU-Net.
4. **Output:** A compressed NIfTI file (.nii.gz) representing the anatomical volume.

3.2.4 Dataset Organization and Structure

Following conversion, the data is restructured according to the rigid nnU-Net raw data format. This structure is inspired by the Medical Segmentation Decathlon (MSD) organization, which facilitates the automated parsing and cross-validation of datasets. Proper adherence to this structure is a prerequisite for the nnUNetv2_plan_and_preprocess command to function.

The directory hierarchy is established as follows:

1. nnUNet_raw/DatasetXXX_Name/: The root folder for the task, where XXX is a unique three-digit identifier (e.g., 501) and Name describes the task (e.g., KidneyAbnormality).

2. `imagesTr/`: Contains the training images (CT scans). Files must follow the naming convention `Case_Identifier_0000.nii.gz`. The `_0000` suffix indicates the modality channel (CT).
3. `labelsTr/`: Contains the corresponding ground truth segmentation masks. Files must be named `Case_Identifier.nii.gz` (without the modality suffix).
4. `dataset.json`: A descriptor file defining the dataset properties. This JSON file acts as a manifest, containing:
 - a. `channel_names`: Mapping of channel IDs to modality types (e.g., `{"0": "CT"}`).
 - b. `labels`: Mapping of integer pixel values to class names (e.g., `{"0": "background", "1": "kidney", "2": "cyst", "3": "stone", "4": "tumor"}`).
 - c. `numTraining`: The total count of training cases.
 - d. `file_ending`: The file extension used (`.nii.gz`).

3.2.5 Dataset Characterization

The initial phase of the nnU-Net pipeline involves the extraction of a “dataset fingerprint,” which captures the intrinsic spatial and intensity properties of the CT data acquired from sources. This fingerprint guides the automated configuration of patch sizes, batch sizes, and network depth. The key characteristics identified during this heuristic analysis are summarized in Table 1.

Property	Value/Description
Modality	Computed Tomography (CT)
Image Size (Voxels)	Variable (Standardized via Resampling)
Target Voxel Spacing	[1.0, 1.0, 1.0]

HU Clipping Range	-135 HU to 215 HU (Automatically set by nnU-Net v2)
Class Distribution	5 Classes: Background, Kidney, Cyst, Stone, Tumor
Data Selection Pool	450 Volumetric Cases (150 cases per entity)
Held-out Test Set	60 Volumetric Cases (20 Per entity)
Final Curated Set	180 Volumetric Cases (60 Per entity)
Data Splitting	5-Fold Cross-Validation

Table 1. Dataset Fingerprint Characteristics

While the number of unique patient cases is 60 per entity, the extraction of all three standard anatomical planes — axial, coronal, and sagittal — yields approximately 300 2D slices per case, resulting in roughly 18,000 slices per pathological class and 54,000 slices across the full training corpus. Furthermore, the patch-based training paradigm of nnU-Net — where the network samples multiple overlapping 3D sub-volumes per case during each training epoch — and the aggressive on-the-fly data augmentation strategy (detailed in 3.3.5) collectively ensure that the effective number of training samples seen by the model over 1,000 epochs substantially exceeds the raw case count. This design mirrors precedents in the medical imaging literature where compact, high-quality volumetric datasets have achieved competitive segmentation performance when paired with self-configuring frameworks such as nnU-Net [51], [58].

3.3 Data Preprocessing

Data preprocessing in medical image analysis is a fundamental step that transforms raw physical measurements into a normalized numerical space suitable for statistical learning. In the context of deep learning, where models are sensitive to the distribution and

scale of input data, preprocessing is not merely a formatting step but a physiological normalization process. The goal is to standardize the physical properties of the input data so that the neural network learns to associate specific numerical values and spatial shapes with anatomical structures, rather than learning artifacts of the acquisition process or inter-patient variability. The preprocessing pipeline for this thesis involves three critical operations: Hounsfield Unit (HU) clipping, isotropic resampling, and Z-score normalization.

3.3.1 Hounsfield Unit (HU) Clipping and Intensity Filtering

Computed Tomography (CT) images provide a quantitative measure of radiodensity, expressed in Hounsfield Units (HU). The HU scale is a linear transformation of the linear attenuation coefficient of the material, calibrated such that the radiodensity of distilled water at standard temperature and pressure (STP) is defined as 0 HU, and the radiodensity of air is defined as -1000 HU. This standardization is a unique advantage of CT over MRI (where intensity values are relative), as it allows for the absolute differentiation of tissues based on their physical density.

However, the full dynamic range of a modern CT scanner can extend from -1000 HU (air) to over +3000 HU (dense metal implants or cortical bone). For the specific task of segmenting kidneys and their associated soft-tissue abnormalities, the vast majority of relevant anatomical information is contained within a much narrower range. "Soft tissues" generally fall between -100 and +200 HU. Feeding the full dynamic range into the network would compress the contrast resolution of

the soft tissues, making it difficult for the model to discern subtle differences between a cyst (fluid) and a tumor (solid tissue).

To assist the neural network in focusing on the relevant anatomical features and to eliminate outliers, intensity clipping is applied.

Mechanism: All voxel values below a defined lower threshold are set to that lower threshold, and all values above a defined upper threshold are set to that upper threshold.

Parameters: The clipping range is explicitly set to -135 HU to 215 HU.

3.3.2 Physiological Justification for Clipping Range:

1. **Lower Bound (-135 HU):** This value is strategically chosen to exclude air (-1000 HU) and the lowest density fat tissues. Perinephric fat (fat surrounding the kidney) typically ranges from -50 to -100 HU. By clipping at -135 HU, the network effectively treats all air (e.g., in bowel loops) and deep adipose tissue as a uniform background value. This reduces the variance the model needs to process in non-relevant regions, effectively "masking out" the dark background and highlighting the kidney within the retroperitoneal fat capsule.
2. **Upper Bound (215 HU):** This value encompasses all relevant soft tissues, including the renal parenchyma (approx. 30-50 HU unenhanced, 100-180 HU contrast-enhanced), cysts (0-20 HU), and solid tumors (variable, typically 40-100 HU). Values significantly higher than 215 HU typically represent cortical bone or highly calcified kidney stones.

While kidney stones are a target class, they are sufficiently dense (>400 HU) that even when clipped at 215 HU, they appear as maximum-intensity objects (bright white) relative to the surrounding soft tissue. Clipping at 215 HU prevents the network's weights from being skewed by the extreme values of metal artifacts or dense bone, while still preserving the high contrast of stones. It focuses the dynamic range on the subtle contrast differences required to distinguish a tumor from a cyst.

Following clipping, the data is normalized. Since CT data has a consistent physical meaning, nnU-Net employs a specific normalization scheme. The clipped values are subjected to a global Z-score normalization:

$$X_{norm} = \frac{x - \mu}{\sigma}$$

where x is the pixel value, μ is the mean intensity of the foreground pixels (computed over the entire training dataset), and σ is the standard deviation. This ensures that the input distribution has a mean of 0 and a unit variance, which is crucial for the convergence stability of gradient descent algorithms during training.

3.3.3 Isotropic Resampling and Spatial Normalization

CT scans are volumetric data composed of stacked 2D slices. A significant challenge in medical imaging is anisotropy, where the physical resolution within a slice (pixel spacing, e.g., 0.7×0.7 mm) is much higher than the resolution between slices (slice thickness, e.g., 3.0 mm or 5.0 mm). This results in voxels that are rectangular prisms rather than cubes. If raw anisotropic data were fed into a 3D CNN, the network would learn distorted spatial filters. For example, a spherical

kidney stone would appear as a flattened pancake in the voxel space if the slice thickness is large. This would severely impair the network's ability to recognize the true 3D morphology of tumors and stones.

To address this, the methodology mandates resampling to a uniform voxel spacing.

Target Spacing: The dataset is resampled to [1.0, 1.0, 1.0] mm isotropic spacing.

Implication: This transformation creates a new voxel grid where every voxel is a perfect cube of 1mm edges.

Algorithm:

1. Image Data: Third-order spline interpolation is used for the CT images. Spline interpolation fits a smooth curve through the surrounding pixel values to calculate the intensity of the new voxels. This preserves the continuous nature of tissue density and maintains the subtle gradients at the boundaries of lesions.
2. Label Masks: Nearest-neighbor interpolation is strictly used for the segmentation masks. Because the masks consist of categorical integers (0, 1, 2, 3, 4), using spline interpolation would introduce fractional values (e.g., 1.5), which are meaningless class identifiers. Nearest neighbor ensures that the class labels remain discrete integers.

This resampling step ensures that physical distances are represented consistently across all three axes (Axial, Sagittal, Coronal), allowing the 3D U-Net to perform spatial convolutions effectively in all directions and recognize the true volumetric shape of renal abnormalities.

After resampling to [1.0, 1.0, 1.0] mm isotropic spacing, all three standard anatomical planes — axial, coronal, and sagittal — are extracted from each volumetric case. Because the voxels are isotropic cubes, all three planes carry equivalent resolution and are geometrically consistent, making multi-planar extraction a straightforward augmentation of the training corpus without introducing resolution artifacts.

3.3.4 Cropping to Non-Zero Region

On datasets that are not pre-processed; to optimize computational efficiency and reduce memory usage, the preprocessing pipeline includes an automated cropping step. Medical images often contain large areas of background air surrounding the patient's body. nnU-Net automatically detects the non-zero region (the body area) and crops the volume to the bounding box of the patient. This significantly reduces the size of the input tensor, allowing for larger batch sizes or patch sizes during training, which contributes to better context modeling.

3.3.5 Augmentation Setting and Data Splitting

Given the high dimensionality of 3D medical data and the limited number of annotated samples available in clinical datasets, deep learning models are prone to overfitting. To mitigate this, a comprehensive data augmentation strategy is combined with a rigorous, automated data splitting scheme.

Data Augmentation Strategy

Category	Transformation	Description and Clinical Justification
Spatial	Rotation & Scaling	Random rotations (z-axis: [-180, 180], x/y: [-30, 30]) and scaling (0.7x - 1.4x) simulate

		variations in patient positioning and differences in organ size.
	Elastic Deformation	Simulates natural soft-tissue deformation caused by respiration or pressure from adjacent organs.
	Mirroring	Random mirroring along axes to leverage the approximate bilateral symmetry of the kidneys.
Intensity	Gaussian Noise/Blur	Simulates scanner noise and reconstruction variability (e.g., soft vs. bone kernels).
	Brightness/Contrast	Makes the model invariant to slight variations in HU calibration or contrast enhancement phases.
	Gamma Correction	Non-linear intensity transformation to simulate diverse tissue density shifts.

Table 2. Data Augmentation Strategy

Data Splitting: Automated 5-Fold Cross-Validation

The data splitting process in this study is handled automatically by the nnU-Net v2 pipeline, which implements a 5-fold cross-validation scheme by default to ensure unbiased performance estimation.

The Automated Splitting Mechanism:

Pipeline-Driven Partitioning: Once the dataset is standardized and the dataset.json is parsed, nnU-Net automatically partitions the total training cases into 5 equal, non-overlapping subsets (folds).

80/20 Distribution: For each of the 5 iterations, the framework uses 4 folds (80% of the data) for model training and the remaining fold (20% of the data) as the internal validation set.

Cross-Validation Iterations: The framework generates 5 independent sets of model weights (Fold 0 through Fold 4). This ensuring that every single case in the dataset is used for validation exactly once, providing a comprehensive assessment of the model's ability to generalize to unseen data.

Held-out Test Set (Manual Allocation)

In addition to the automated 5-fold internal cross-validation, a subset of the dataset is manually reserved as an independent held-out test set.

Purpose: This set is never seen by the model during the training or validation iterations of the 5-fold process. It serves as the final arbiter of performance, simulating how the AI would perform on entirely new patients in a clinical deployment scenario.

Ensemble Strategy: For the final evaluation on this test set, an ensemble of the **5-fold** models is used. The softmax probability maps from all 5 models are averaged voxel-wise to generate the final mask, a technique that typically boosts accuracy and robustness.

3.4 Network Architecture: The 3D U-Net and Deep Supervision

The core computational engine of this methodology is the 3D U-Net, a fully convolutional neural network specifically designed for volumetric segmentation. While the

U-Net architecture is a standard in biomedical imaging, the nnU-Net v2 implementation introduces dynamic adaptation and deep supervision to maximize performance.




Figure 17. nnU-Net V2 Architecture

3.4.1 Architectural Topology

The network follows the classic symmetric Encoder-Decoder structure with skip connections, adapted for 3D inputs:

1. The Encoder (Contracting Path): This pathway functions as a feature extractor. It consists of repeated blocks of convolutional layers. In nnU-Net v2, these blocks typically comprise:

3D Convolution: 3 x 3 x 3 kernels to extract spatial features.

Instance Normalization: Used instead of Batch Normalization because batch sizes in 3D training are often small (due to memory constraints), rendering Batch Norm unstable. Instance Norm normalizes each sample independently, preserving the contrast information of the individual scan.

Leaky ReLU: A non-linear activation function $f(x) = \max(0.01x, x)$ that introduces non-linearity while allowing a small gradient to flow even for negative values, preventing the "dying ReLU" problem.

Strided Convolution: Used for downsampling (reducing spatial resolution) instead of max-pooling. This allows the network to learn the optimal downsampling operation.

2. The Bottleneck: The deepest part of the network represents the input volume in a highly abstract, low-resolution feature space, capturing semantic context (e.g., "presence of a tumor") rather than precise localization.
3. The Decoder (Expanding Path): This pathway recovers the spatial resolution. It utilizes transposed convolutions to upsample the feature maps.

Skip Connections: Critical to the U-Net design, these connections concatenate the high-resolution feature maps from the Encoder directly to the corresponding layers in the Decoder. This re-injects the fine-grained spatial details (gradients, edges) lost during down sampling, enabling the decoder to precisely localize the boundaries of small structures like kidney stones.

3.4.2 Deep Supervision

A key conceptual element of the nnU-Net methodology is Deep Supervision. In standard networks, the loss is calculated only at the final output layer. However, in deep 3D networks, gradients can vanish as they backpropagate through many layers, making the training of the early layers difficult.

nnU-Net addresses this by outputting segmentation maps not just at the final full-resolution layer, but also at lower-resolution stages of the decoder.

1. **Mechanism:** Auxiliary segmentation heads are attached to the lower-resolution blocks of the decoder.
2. **Loss Calculation:** The loss is computed for each of these outputs (compared against down sampled versions of the ground truth).
3. **Weighted Sum:** The total loss is a weighted sum of the losses at different resolutions. The weights typically decay exponentially (e.g., $1/2$, $1/4$, $1/8$) for lower resolutions, prioritizing the final high-resolution output.
4. **Benefit:** This forces the deeper layers of the network to learn semantically meaningful representations early in the training process and improves

gradient flow, leading to faster and more stable convergence. It ensures that the network pays attention to both coarse context (organ localization) and fine detail (lesion boundaries).

3.5 Computational Pipeline

The training pipeline is meticulously configured to ensure convergence in the high-dimensional non-convex optimization landscape of 3D segmentation.

3.5.1 Environment and Reproducibility

To ensure the computational reproducibility of this study, all experiments were conducted within a standardized software and hardware environment. The model will be trained using Python 3.11 and the PyTorch deep learning framework, utilizing nnU-Net v2. High-intensity computations were executed on a single NVIDIA A100 GPU with 40GB VRAM, hosted via Google Colab Pro+. This hardware configuration provided the necessary VRAM budget to support large patch sizes and optimized batch gradients, ensuring stable convergence during the 1000-epoch training cycle.

Component	Technical Detail
Deep Learning Framework	PyTorch 2.x / nnU-NetV2
Programming Language	Python 3.11
GPU Hardware	NVIDIA A100 (40GB VRAM)
Computing Platform	Google Colab Pro+
Conversion Utility	dcm2niix (Python Library)
Annotation Platform	3D Slicer

Table 3. Technical Specifications*3.5.2 Training Configuration and Optimization*

Optimizer: The Stochastic Gradient Descent (SGD) optimizer is selected, utilizing Nesterov momentum. While adaptive optimizers like Adam are popular, empirical evidence from the nnU-Net authors suggests that SGD with Nesterov momentum generalizes better for medical segmentation tasks.

Hyperparameter	Value/Configuration
Training Schedule	1000 Epochs (250 iterations per epoch)
Optimization Algorithm	SGD with Nesterov Momentum (0.99)
Initial Learning Rate	10^{-2} (Polynomial decay schedule)
Weight Decay	3×10^{-5}
Objective Function	Dice Similarity + Cross – Entropy Loss
Validation Strategy	Automated 5-fold Cross-Validation

Table 4. Training Hyperparameters

Momentum: Set to 0.99. This high momentum helps the optimizer navigate the complex loss landscape, allowing it to barrel through shallow local minima and converge to a flatter, more robust minimum.

Weight Decay: Set to 3×10^{-5} to prevent overfitting by penalizing large weights (L2 regularization).

3.5.3 Loss Function Math

The training objective is minimized using a compound loss function that combines Dice Loss and Cross-Entropy (CE) Loss:

$$L_{total} = L_{Dice} + L_{CrossEntropy}$$

1. Dice Loss: Directly optimizes the metric of interest (Dice score). It is robust to class imbalance because it relies on the intersection-over-union concept.

Even if the background is 99% of the image, the Dice loss focuses solely on the overlap of the foreground classes.

2. Cross-Entropy Loss: Provides dense, pixel-wise supervision by penalizing incorrect class probability assignments for every voxel in the volume. Unlike Dice Loss, which is computed globally over the segmentation mask, CE Loss enforces local correctness at each voxel, improving boundary sharpness and stabilizing training in the early epochs when Dice gradients are weak.

Training Schedule:

1. Epochs: The model is trained for a fixed duration of 1000 epochs. Uniquely, an "epoch" in nnU-Net is defined not by the dataset size, but as a fixed number of iterations (250 mini-batches). This results in a total of 250,000 training steps.
2. Learning Rate (LR): A "poly" learning rate policy is employed. The initial LR (typically 0.01) is decayed following a polynomial curve:

$$LR_{iter} = LR_{initial} \times \left(1 - \frac{iter}{maxiter}\right)^{0.9}$$

This allows for rapid exploration of the loss landscape in the early phases and precise fine-tuning in the later phases.

3.5.4 Inference and Post-Processing

Following the inference phase, nnU-Net v2 evaluates the validity of the predicted masks through automated post-processing. This primarily involves Connected Component Analysis (CCA), where the framework identifies and removes small, isolated clusters of pixels that are statistically unlikely to represent

actual anatomical structures or lesions. By eliminating these "islands" of noise, the methodology ensures the generation of more coherent and clinically plausible segmentation masks, which is particularly vital for accurately delineating high-contrast structures like kidney stones in complex abdominal scans.

3.6 Annotation Pipeline

The validity of any supervised learning model is bounded by the quality of its ground truth. We employed a dual-stream data ingestion strategy:

- a. **Pre-annotated Data:** For the KiTS (Kidney Tumor Segmentation) repository, which includes expert-verified segmentation masks, we utilized the provided ground truth labels. These masks underwent a Verification Protocol to ensure compatibility with our class mapping and NIfTI conversion standards.
- b. **Raw Data:** For supplementary datasets (e.g., specific stone or cyst cases collected without segmentation), we implemented a Manual Annotation Pipeline to generate high-precision, radiologist-verified masks.

Tool: 3D Slicer, an open-source software platform for medical image informatics and visualization, is utilized for all manual segmentation tasks. The "Segment Editor" module provides advanced tools for painting, drawing, and refining 3D volumes.

Labeling Protocol: Each voxel in the CT volume is assigned a unique integer label representing its tissue class:

Integer Label	Clinical Entity	Description
0	Background	All non-renal tissues and air
1	Kidney	Healthy renal parenchyma

2	Cyst	Fluid-filled renal lesions
3	Stone	High-density calcifications (>400 HU)
4	Tumor	Solid renal neoplasms

Table 5. Class Mapping

3.6.1 Domain Expert Profile

To ensure the clinical validity of the dataset, the ground truth annotations and failure analysis were supervised by Dr. Nathaniel F. Paragas. As an Affiliate Consultant at Capitol Medical Center—and an Associate Professor at Mapúa University, Dr. Paragas provides the necessary expert oversight. His specialization in diagnostic and interventional radiology ensures the anatomical precision required for differentiating complex renal cysts, tumors, and stones.

Category	Details
Name & Title	Dr. Nathaniel F. Paragas, MD, FPCR, FPSVIR
Current Affiliations	Associate Professor, School of Medicine, Mapúa University
	Affiliate Consultant, Department of Radiology, Capitol Medical Center
Specialization	Diagnostic Radiology (Residency, St. Luke's Medical Center
	Vascular & Interventional Radiology (Fellowship, St. Luke's
Experience	10 Years – Medical Doctor
	6 Years – Diagnostic Radiology
	3 Years – Vascular & Interventional Radiology
Professional Status	Fellow, Philippine College of Radiology (FPCR)

	Fellow, Philippine Society of Vascular and Interventional Radiology (FPSVIR)
--	--

Table 6. Domain Expert (Radiologist) Profile

To supplement the radiological expertise in image annotation and model validation, a second domain expert with nephrology specialization is included. This nephrologist provides clinical insight into the pathophysiological presentation of renal abnormalities, particularly regarding the differentiation of complex cystic lesions and the clinical significance of stone burden measurements. The inclusion of a nephrologist ensures that failure analysis and boundary validation are evaluated from both an imaging-technical and a disease-management perspective.

Workflow:

1. Initialization: The NIfTI volume is loaded into 3D Slicer. A new segmentation layer is created.
2. Label Generation:
 - a) Stream A (Raw Data): Researchers manually 'painted' regions of interest using 3D Slicer's interpolation tools, followed by slice-by-slice correction.
 - b) Stream B (Pre-annotated Data): Existing masks were mapped to the study's label values (e.g., ensuring KiTS 'Tumor' label maps to Class ID 4) and visually inspected for conversion artifacts."
3. Expert Review: A "Human-in-the-Loop" quality assurance process is implemented.

A. For Raw Data: Dr. Paragas supervised the manual delineation process, reviewing 100% of the student-generated masks for complex cysts and stones to validate boundary precision.

a) Pilot Phase: A subset of 100 cases is segmented and reviewed by the collaborating radiologist expert (Dr. Nathaniel Paragas). The radiologist provides feedback on the distinction between complex cysts and tumors, and on the precise boundaries of lesions.

b) Refinement: Based on this feedback, the annotations are refined.

c) Full-Scale: The approved protocol is applied to the entire dataset. Consistency checks are performed to identify and correct any missed lesions or mislabeled regions.

B. For Pre-annotated Data (KiTS): A random sampling of 10% of cases was reviewed to confirm that the imported labels-maintained integrity during the DICOM-to-NIfTI conversion and resolution resampling.

The radiologist expert does not manually segment every slice from scratch, as this is time prohibitive. Instead, a corrective annotation workflow is employed: student researchers generate the initial masks, and the collaborating radiologist reviews and refines the 3D volumes using 3D Slicer's verification tools, focusing on complex boundaries and ambiguous lesions.

3.7 Evaluation Metrics

To quantitatively assess the performance of the segmentation model, a comprehensive set of metrics is utilized. These metrics are selected to capture complementary aspects of segmentation quality: volumetric overlap and boundary precision.

3.7.1 Dice Similarity Coefficient (DSC)

The Dice Similarity Coefficient (DSC) (also known as the Sørensen–Dice index) is the standard metric for validating medical image segmentation. It measures the spatial overlap between the automated segmentation (A) and the manual ground truth (B).

$$DSC = \frac{2 |A \cap B|}{|A| + |B|} = \frac{2 * TP}{2 * TP + FP + FN}$$

where TP is True Positives (voxels correctly classified as the target), FP is False Positives (background voxels classified as target), and FN is False Negatives (target voxels classified as background).

Clinical Relevance: The DSC ranges from 0 (no overlap) to 1 (perfect overlap). In clinical contexts like radiotherapy planning or tumor burden monitoring, the total volume of the lesion is critical. A high DSC ensures that the AI-predicted volume is statistically consistent with the expert's measurement.

3.7.2 Intersection over Union (IoU)

The Intersection over Union, also known as the Jaccard Index, computes the ratio of the intersection to the union of the predicted and ground truth segmentation masks. IoU is a stricter overlap metric

than DSC because the union term more heavily penalizes both over-segmentation and under-segmentation simultaneously.

$$\text{IoU} = |\mathbf{A} \cap \mathbf{B}| / |\mathbf{A} \cup \mathbf{B}| = \text{TP} / (\text{TP} + \text{FP} + \text{FN})$$

where A is the predicted segmentation mask and B is the ground truth mask.

The relationship between DSC and IoU is defined as:

$$\text{DSC} = (2 \times \text{IoU}) / (1 + \text{IoU})$$

This relationship implies that IoU will always yield a lower numerical value than DSC for the same prediction. This distinction is important when interpreting and comparing results across studies that may report only one of the two metrics.

3.7.3 95th Percentile Hausdorff Distance (HD95)

While DSC is excellent for measuring volume, it is relatively insensitive to shape errors, especially for large objects. A segmentation could have a high DSC but still contain a significant protrusion or "spike" that extends into a neighboring organ. In surgery, such an error could be catastrophic.

To measure boundary accuracy, the Hausdorff Distance is used. It calculates the maximum distance from a point in one set to the nearest point in the other set. Specifically, the 95th Percentile Hausdorff Distance (HD95) is employed.

Formula:

$$HD(A, B) = \max (\sup_{a \in A} \inf_{b \in B} d(a, b), \sup_{b \in B} \inf_{a \in A} d(b, a))$$

3.7.4 Precision, Recall, and F1-Score

$$Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 * Precision * Recall}{Precision + Recall} = \frac{2 * TP}{2 * TP + FP + FN}$$

To provide a comprehensive assessment of segmentation quality, three additional metrics are reported per class. Precision (Positive Predictive Value) measures the proportion of predicted positive voxels that are truly positive, quantifying the rate of false alarms. Recall (Sensitivity) measures the proportion of actual positive voxels correctly identified by the model, capturing the completeness of lesion detection. F1-Score is the harmonic mean of precision and recall, providing a single balanced metric that penalizes both over-segmentation and under-segmentation equally. These metrics are reported per class (cyst, stone, tumor) for all four model configurations to enable direct cross-model comparison.

3.7.5 Per-Class Analysis

A single global metric can be misleading. The kidney is a large structure, while a stone is tiny. A model could achieve a 0.90 global

Dice score by segmenting the kidney perfectly while missing every single stone. Therefore, the evaluation will report DSC and HD95 separately for each class (Kidney, Cyst, Stone, Tumor). This granular analysis is essential to identify specific failure modes—for example, if the model struggles specifically with differentiating dense stones from contrast-enhanced vessels.

3.7.6 Computational Efficiency Metrics

To assess the clinical feasibility of integrating the proposed model into real-time radiological workflows, computational efficiency is evaluated alongside accuracy.

Inference Latency: The average time required to process a single patient volume (from raw NIfTI input to final segmentation output) is recorded. This is measured in seconds per volume (s/vol) across the validation set.

Resource Utilization: The peak GPU Video Random Access Memory (VRAM) consumption during inference is monitored to determine the minimum hardware requirements for deployment.

These metrics are recorded using the native PyTorch CUDA event timer and memory manager to ensure precise benchmarking independent of system overhead.

3.7.7 Baseline Benchmarks

To rigorously evaluate whether a unified multi-class nnU-Net V2 architecture performs equivalently or superior to specialized single-

class models, strict quantitative baselines are established from contemporary literature using identical or closely related architectures (3D U-Net and nnU-Net). Table 7 aggregates benchmark metrics from specialized models published between 2022 and 2025.

Reference	Target Pathology	Architecture	DSC	IoU	Precision	Recall
Elton et al. [24]	Kidney Stones (NCCT)	Custom CNN	-	-	0.966	0.978
Li et al. [39]	Kidney / Stones	3D U-Net	0.960	-	-	0.802 (Stones)
Soleimany et al. [60]	Renal Tumors (KiTS19)	nnU-Net / 3D U-Net	0.960 (Tumor)	0.910	-	-
Doostinia et al. [63]	Renal Masses (KiTS23)	nnU-Net V2	0.645 (Tumor)	-	-	-
de Boer et al. [65]	Cysts & Tumors Combined	nnU-Net	0.660 (Abnormality)	-	-	-

Table 7. Baseline Metrics for Kidney Abnormality Segmentation

Note: Missing values indicate the metric was not the primary optimization target in that study. Elton et al. [24] prioritized detection precision/recall due to the minute volumetric size of stones. Critically, no published benchmark provides DSC or HD95 for kidney stone segmentation under nnU-Net, as the major challenge datasets (KiTS19, KiTS23) do not include stones as a segmentation class. The stone-class results from this study's single-class and multi-

class nnU-Net V2 configurations will constitute the first reported stone segmentation benchmarks under the nnU-Net V2 framework.

The 3D U-Net baseline in this study must approximate Li et al.'s [39] 0.802 sensitivity for kidney stones and approach Soleimany et al.'s [60] 0.960 DSC for tumor boundaries before the Phase 2 multi-class comparison is interpreted.

3.8 Failure Analysis and Clinical Interpretability

To bridge the gap between computational metrics and clinical applicability, a qualitative Failure Analysis phase is integrated into the methodology.

Protocol:

1. Identification: The validation cases with the lowest Dice scores and highest HD95 values are extracted from the cross-validation results.
2. Visual Inspection: These "failure cases" are loaded into 3D Slicer, with the AI prediction overlaid on the ground truth and the original CT scan.
3. Radiologist Review: The radiologist expert reviews these cases to determine the root cause of the failure.
4. Categorization: Errors are classified into categories such as:
 - a) Image Artifacts: Failure due to metal streak artifacts or low signal-to-noise ratio.
 - b) Ambiguous Anatomy: Cases where even the human expert is unsure of the boundary (e.g., a cyst with indeterminate density).
 - c) Model Hallucinations: False positives where the model detects a lesion that does not exist (e.g., misclassifying a vessel as a stone).

- d) Boundary Leakage:** The model correctly detects the tumor but extends the segmentation into adjacent muscle or fat.

This qualitative analysis provides the "why" behind the numbers, offering crucial insights for future model improvements and establishing realistic expectations for the tool's clinical reliability.

3.9 Clinical Feature Quantification

While deep learning models output voxel-wise segmentation masks, clinical decision-making requires quantitative metrics. This study implements a deterministic post-processing algorithm to convert raw segmentation masks into clinical descriptors necessary for Radius Exophytic/Endophytic Nearness Anterior/Posterior Location (R.E.N.A.L.) Nephrometry scoring and stone burden analysis

3.9.1 Volumetric Calculation

The volume (V) of any segmented structure (tumor, cyst, or stone) is calculated by counting the total number of positive voxels (N_{voxels}) belonging to that class and multiplying by the volume of a single voxel (V_{spacing}), derived from the resampled isotropic spacing (1.0mm x 1.0mm x 1.0mm):

$$V = N_{\text{voxels}} \times V_{\text{spacing}}$$

3.9.2 Maximum Diameter (Tumor Size)

To determine the maximal diameter required for the 'R' (Radius) component of nephrometry scoring, we utilize Principal Component Analysis (PCA) on the voxel coordinates of the tumor mask to find the major axis length, or compute the maximum Euclidean distance between any two surface voxels:

$$D_{max} = \max_{\{i,j \in S\}} \|p_i - p_j\|_2$$

3.10 Testing and Validation Pipeline

This study executes exactly four nnU-Net V2 training sessions organized into two sequential phases. Phase 1 trains three single-class nnU-Net V2 models — one per abnormality class (cyst, stone, tumor) — and validates each against published benchmarks. Phase 2 trains one multi-class nnU-Net V2 model on all three classes simultaneously and compares its per-class performance against the Phase 1 models. A standardized testing pipeline is applied identically across all configurations.

3.10.1 Experimental Configuration Matrix

The four training sessions are structured as follows:

Phase 1 — Single-Class Validation (Three Training Sessions)

1. Configuration A: Single-class nnU-Net V2 trained on stones only. Cysts and tumors are treated as background.
2. Configuration B: Single-class nnU-Net V2 trained on cysts only.
3. Configuration C: Single-class nnU-Net V2 trained on tumors only.

Phase 1 outputs per-class DSC, IoU, HD95, Precision, Recall, and F1-Score for each dedicated model. These are benchmarked against the literature metrics in Table 7 (Section 3.7.7). This validation is a prerequisite — a single-class model that significantly underperforms published benchmarks is a weak reference, and any multi-class advantage observed in Phase 2 could be falsely attributed to baseline weakness rather than genuine multi-class learning benefit.

Phase 2 — Multi-Class Comparison (One Training Session)

1. Configuration D: Multi-class nnU-Net V2 trained simultaneously on all three classes, producing a voxel-wise softmax probability map over background, kidney, cyst, stone, and tumor.

Configuration D is the primary focus of this study. It is evaluated on the same held-out test set under identical preprocessing and inference conditions as Phase 1, and its per-class metrics are compared directly against Configurations A, B, and C.

3.10.2 Held-Out Test Set Preparation

The 60 volumetric cases reserved as the held-out test set (20 per pathological class) are never exposed to any model during training or cross-validation. Prior to inference, each test volume undergoes the same preprocessing pipeline applied during training: DICOM-to-NIfTI conversion, HU intensity clipping to the range of -135 to 215 HU, isotropic resampling to 1.0 mm spacing, and Z-score normalization. This ensures that all test inputs exist in an identical numerical space to the training data, eliminating preprocessing as a source of performance variance across configurations.

3.10.3 Inference & Post-Processing

Each preprocessed test volume is passed through each of the four model configurations independently. For all nnU-Net V2 configurations, the ensemble of all five cross-validation fold models is used during inference. The softmax

probability maps from each of the five folds are averaged voxel-wise to produce the final segmentation mask. This ensemble strategy reduces prediction variance and produces more stable outputs compared to any single fold model. The output of each inference pass is a voxel-wise segmentation mask assigning each voxel to one of the target classes: background, cyst, stone, or tumor.

The raw predicted masks are subjected to Connected Component Analysis (CCA) to remove small, isolated voxel clusters that are statistically unlikely to represent true anatomical structures. Any predicted region falling below a minimum voxel count threshold is discarded as noise. This step ensures that all evaluated masks are clinically plausible and free of spurious detections before metric computation is performed.

3.10.4 Metric Computation

For each test case and each model configuration, the following metrics are computed per class by comparing the predicted segmentation mask against the ground truth annotation:

Metric	Purpose
Dice Similarity Coefficient (DSC)	Volumetric overlap between prediction and ground truth
Intersection over Union (IoU)	Strict overlap ratio penalizing both over- and under-segmentation
95 th Percentile Hausdorff Distance (HD95)	Worst-case boundary deviation
Precision	Proportion of predicted positive voxels that are truly positive
Recall	Proportion of true positive voxels correctly captured by the model
F1-Score	Harmonic balance of precision and recall

Table 8. Metric Computation*3.10.5 Computational Efficiency Recording*

Alongside segmentation accuracy metrics, the following computational efficiency measures are recorded for each model configuration during inference on the test set:

1. Inference Latency: the average time in seconds required to process a single patient volume from preprocessed NIfTI input to final post-processed segmentation output, measured using PyTorch CUDA event timers
2. Peak GPU VRAM Usage: the maximum video memory consumed during a single inference pass, monitored via PyTorch’s native memory manager

These values are averaged across all test cases per configuration to produce a representative efficiency profile, enabling direct comparison of deployment feasibility between the single-class and multi-class approaches.

3.10.6 Cross Configuration Comparison

The computed metrics are organized into a structured comparison table across all four configurations per class. The comparison proceeds in two tiers mirroring the experimental design. First, Phase 1 single-class model results are compared against the published benchmarks in Table 7 (Section 3.7.7) to confirm baseline validity. Second — and as the primary analysis — the multi-class Configuration D results are compared per-class against the three single-class Configurations A, B, and C to determine whether joint training is beneficial, neutral, or detrimental for each pathology. To assess statistical significance, a

paired Wilcoxon signed-rank test is applied to DSC, IoU, and HD95 scores across all test volumes.

3.10.7 Domain Expert Review

The predicted segmentation masks from the held-out test set are reviewed by the two domain experts — the diagnostic radiologist and the nephrologist — with review focused on cases flagged by low DSC scores or high HD95 values. This qualitative review serves to validate whether computational failures correspond to clinically meaningful errors or whether they represent boundary ambiguities that are inherently difficult even for human annotators. Expert findings are recorded and fed directly into the Failure Analysis phase described in Section 3.9.

3.11 Web Deployment

To bridge the translational gap between the computational backend and clinical end-users (radiologists and nephrologists), this study implements a web-based deployment interface. This component serves as the delivery mechanism for the Input-Process-Output model, transforming the trained nnU-Net v2 model from a command-line research tool into an accessible diagnostic prototype.

3.11.1 Framework Selection: Streamlit

The user interface is developed using Streamlit, an open-source Python framework designed specifically for machine learning and data science applications. Streamlit is selected for this study due to its seamless integration with the PyTorch backend, allowing for the direct execution of inference scripts without

the latency overhead of intermediate API calls associated with traditional web frameworks (e.g., Django or Flask). Its capability to render interactive widgets and visualize multidimensional arrays (3D medical volumes) directly from Python scripts ensures rapid prototyping and maintains the fidelity of the segmentation results.

3.11.2 Application Architecture

The deployment architecture follows a three-tier design optimized for clinical workflow simulation:

1. **Data Ingestion Layer:** The interface provides a secure "Drag-and-Drop" uploader supporting both raw DICOM series and pre-converted NIfTI files. For DICOM inputs, the application triggers an integrated `dcm2niix` background process to standardize the volume before inference, mirroring the Input Phase described in Section 3.1.2.
2. **Inference Engine:** Upon upload, the application instantiates the trained nnU-Net v2 model (specifically the configuration yielding the highest validation dice scores). The uploaded volume is preprocessed (HU clipping and resampling) in real-time, passed through the 3D U-Net, and the resulting segmentation masks are post-processed to remove artifacts.
3. **Visualization Layer:** The application utilizes interactive slice navigators to display the original CT volume alongside the semi-transparent

segmentation overlay. This allows clinicians to scroll through Axial, Coronal, and Sagittal planes to visually verify the alignment of the predicted masks (Kidney, Cyst, Stone, Tumor) with the anatomical structures.

3.11.3 Clinical Reporting Features

Beyond visualization, the Streamlit deployment automates the generation of the quantitative outputs defined in Section 3.9. The interface dynamically computes and displays:

- a) **Volumetric Analysis:** A summary table listing the total volume (in mL) for every detected abnormality class.

- b) **Downloadable Reports:** The system compiles these metrics and representative key images into a standardized PDF report, simulating the integration of AI findings into a radiologist's final diagnostic report.

REFERENCES

- [1] [1] B. Mensel et al., "Prevalence of Renal Cysts and Association with Risk Factors in a General population: an MRI-based Study," *Abdominal Radiology (New York)*, vol. 43, no. 11, pp. 3068-3074, Nov. 2018, doi: <https://doi.org/10.1007/s00261-018-1565-5>.
- [2] B. H. Chew, L. E. Miller, B. Eisner, S. Bhattacharyya, and N. Bhojani, "Prevalence, Incidence, and Determinants of Kidney Stones in a Nationally Representative Sample of US Adults," *JU open plus*, vol. 2, no. 1, Jan. 2024, doi: <https://doi.org/10.1097/ju9.0000000000000107>.
- [3] A. F. Awedew et al., "The global, regional, and national burden of urolithiasis in 204 countries and territories, 2000-2021: a systematic analysis for the Global Burden of Disease Study 2021," *EClinicalMedicine*, vol. 78, pp. 102924-102924, Nov. 2024, doi: <https://doi.org/10.1016/j.eclinm.2024.102924>.

- [4] A. Volpe et al., "Contemporary Management of Small Renal Masses," *European Urology*, vol. 60, no. 3, pp. 501-515, Sep. 2011, doi: <https://doi.org/10.1016/j.eururo.2011.05.044>.
- [5] S. C. Campbell, P. E. Clark, S. S. Chang, J. A. Karam, L. Souter, and R. G. Uzzo, "Renal Mass and Localized Renal Cancer: Evaluation, Management, and Follow-Up: AUA Guideline: Part I," *Journal of Urology*, vol. 206, no. 2, pp. 199-208, Aug. 2021, doi: <https://doi.org/10.1097/ju.0000000000001911>.
- [6] M. Akram and B. Somani, "Epidemiology and Management of Kidney Stone Disease - Current Insights," *Research and Reports in Urology*, vol. 17, pp. 449-459, Nov. 2025, doi: <https://doi.org/10.2147/rru.s517758>.
- [7] S. G. Silverman et al., "Bosniak Classification of Cystic Renal Masses, Version 2019: An Update Proposal and Needs Assessment," *Radiology*, vol. 292, no. 2, pp. 475-488, Aug. 2019, doi: <https://doi.org/10.1148/radiol.2019182646>.
- [8] E. Salles-Silva, E. M. Lima, V. B. Amorim, M. Milito, and D. B. Parente, "Clear Cell Likelihood Score May Improve Diagnosis and Management of Renal Masses," *Abdominal Radiology*, Jun. 2024, doi: <https://doi.org/10.1007/s00261-024-04415-4>.
- [9] N. Ashraf et al., "Incidence and factors associated with burnout in radiologists: A systematic review," *European Journal of Radiology Open*, vol. 11, p. 100530, Oct. 2023, doi: [10.1016/j.ejro.2023.100530](https://doi.org/10.1016/j.ejro.2023.100530).
- [10] J. A. Harolds, J. R. Parikh, E. I. Bluth, S. C. Dutton, and M. P. Recht, "Burnout of Radiologists: Frequency, Risk Factors, and Remedies: a Report of the ACR Commission on Human Resources," *Journal of the American College of Radiology: JACR*, vol. 13, no. 4, pp. 411-416, Apr. 2016, doi: <https://doi.org/10.1016/j.jacr.2015.11.003>.

- [11] G. Litjens et al., "A survey on deep learning in medical image analysis," *Medical Image Analysis*, vol. 42, pp. 60-88, Jul. 2017, doi: 10.1016/j.media.2017.07.005.
- [12] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," *arXiv.org*, May 18, 2015.
<https://arxiv.org/abs/1505.04597>
- [13] P. Mukherjee, S. Lee, D. C. Elton, S. Y. Nakada, P. J. Pickhardt, and R. M. Summers, "Fully Automated Longitudinal Assessment of Renal Stone Burden on Serial CT Imaging Using Deep Learning," *Journal of Endourology*, vol. 37, no. 8, pp. 948-955, Jun. 2023, doi: 10.1089/end.2023.0066.
- [14] M. Zhang et al., "Imaging-based deep learning in kidney diseases: recent progress and future prospects," *Insights Into Imaging*, vol. 15, no. 1, Feb. 2024, doi: 10.1186/s13244-024-01636-5.
- [15] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks," *arXiv.org*, Jun. 04, 2015.
<https://arxiv.org/abs/1506.01497>
- [16] K. A. Abdimurotovich and Y.-I. Cho, "Optimized YOLOv5 Architecture for Superior Kidney Stone Detection in CT Scans," *Electronics*, vol. 13, no. 22, p. 4418, Nov. 2024, doi: 10.3390/electronics13224418.
- [17] T. Zhao, Z. Sun, Y. Guo, Y. Sun, Y. Zhang, and X. Wang, "Automatic renal mass segmentation and classification on CT images based on 3D U-Net and ResNet algorithms," *Frontiers in Oncology*, vol. 13, p. 1169922, May 2023, doi: 10.3389/fonc.2023.1169922.

- [18] A. G. Raman et al., "Evaluation of nnU-Net for kidney tumor segmentation on a large external patient cohort," *European Journal of Radiology Artificial Intelligence*, vol. 3, p. 100035, Jul. 2025, doi: <https://doi.org/10.1016/j.ejrai.2025.100035>.
- [19] N. Heller et al., "The state of the art in kidney and kidney tumor segmentation in contrast-enhanced CT imaging: Results of the KiTS19 challenge," *Medical Image Analysis*, vol. 67, p. 101821, Oct. 2020, doi: [10.1016/j.media.2020.101821](https://doi.org/10.1016/j.media.2020.101821).
- [20] Z. Wu et al., "MSWAL: 3D Multi-class Segmentation of Whole Abdominal Lesions Dataset," *arXiv.org*, 2025. <https://arxiv.org/abs/2503.13560>
- [21] S. Battiato, G. M. Farinella, G. Gallo, O. Garretto, and C. Privitera, "Objective analysis of simple kidney cysts from CT images," in *2009 IEEE Int. Workshop Med. Meas. Appl.*, 2009, pp. 146–149, doi: [10.1109/MEMEA.2009.5167973](https://doi.org/10.1109/MEMEA.2009.5167973)
- [22] R. Reinhard, M. van der Zon-Conijn, and R. Smithuis, "The Radiology Assistant : Solid masses," *radiologyassistant.nl*, Jun. 14, 2016. <https://radiologyassistant.nl/abdomen/kidney/solid-masses>
- [23] A. Shteynshlyuger, "Diagnosis of Kidney Stones and Ureteral Stones - NYC Kidney Stone Clinic," *New York Urology Specialists*. <https://www.newyorkurologyspecialists.com/kidney-stones/diagnosis/>
- [24] D. C. Elton, E. B. Turkbey, P. J. Pickhardt, and R. M. Summers, "A deep learning system for automated kidney stone detection and volumetric segmentation on noncontrast CT scans," *Medical Physics*, vol. 49, no. 4, Feb. 2022, doi: <https://doi.org/10.1002/mp.15518>.
- [25] M. A. Bosniak, "The Bosniak Renal Cyst Classification: 25 Years Later," *Radiology*, vol. 262, no. 3, pp. 781-785, Mar. 2012, doi: <https://doi.org/10.1148/radiol.11111595>.

- [26] F. Gaillard, T. Walizai, R. Thibodeau, et al., "Bosniak classification system of renal cystic masses (version 2005)," Radiopaedia.org. [Online]. Available: <https://doi.org/10.53347/rID-1006>.
- [27] S. Ananth, "Faster R-CNN for object detection," Medium, Aug. 09, 2019. <https://medium.com/data-science/faster-r-cnn-for-object-detection-a-technical-summary-474c5b857b46>
- [28] Md. E. Rayed, S. M. S. Islam, S. I. Niha, J. R. Jim, M. M. Kabir, and M. F. Mridha, "Deep learning for medical image segmentation: State-of-the-art advancements and challenges," *Informatics in Medicine Unlocked*, vol. 47, p. 101504, Jan. 2024, doi: 10.1016/j.imu.2024.101504.
- [29] A. Kabil, G. Khoriba, M. Yousef, and E. A. Rashed, "Advances in medical image segmentation: A comprehensive survey with a focus on lumbar spine applications," *Computers in Biology and Medicine*, vol. 198, no. Pt A, p. 111171, Oct. 2025, doi: 10.1016/j.compbimed.2025.111171.
- [30] P.-H. Conze, G. Andrade-Miranda, Y. L. Meur, E. C.-L. Gall, and F. Rousseau, "Dual task kidney MR segmentation with transformers in autosomal-dominant polycystic kidney disease," *Computerized Medical Imaging and Graphics*, vol. 113, p. 102349, Feb. 2024, doi: 10.1016/j.compmedimag.2024.102349.
- [31] N. Carion, F. Massa, G. Synnaeve, A. Kirillov, and S. Zagoruyko, "End-to-End Object Detection with Transformers," *arXiv.org*, May 26, 2020. <https://arxiv.org/abs/2005.12872>
- [32] A. Dosovitskiy et al., "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale," *arXiv.org*, Oct. 22, 2020. <https://arxiv.org/abs/2010.11929>

- [33] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, You Only Look Once: Unified, Real-Time Object Detection. IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016, pp. 779-788. doi: 10.1109/cvpr.2016.91.
- [34] A. Ramadhani and A. Salam, "Deployment of Web-Based YOLO for CT Scan Kidney Stone Detection," SinkrOn, vol. 8, no. 3, pp. 1357-1368, Jul. 2024, doi: 10.33395/sinkron.v8i3.13744.
- [36] J. Solawetz, "YOLOv7 - A breakdown of how it works," Roboflow Blog, Jul. 17, 2022. <https://blog.roboflow.com/yolov7-breakdown/>
- [37] Y. Yin, Z. Tang, and H. Weng, "Application of visual transformer in renal image analysis," BioMedical Engineering OnLine, vol. 23, no. 1, p. 27, Mar. 2024, doi: 10.1186/s12938-024-01209-z.
- [38] GeeksforGeeks, "Vision Transformer (ViT) Architecture," GeeksforGeeks, Oct. 09, 2024. <https://www.geeksforgeeks.org/deep-learning/vision-transformer-vit-architecture/>
- [39] D. Li et al., "Deep Segmentation Networks for Segmenting Kidneys and Detecting Kidney Stones in Unenhanced Abdominal CT Images," Diagnostics, vol. 12, no. 8, p. 1788, Jul. 2022, doi: 10.3390/diagnostics12081788.
- [40] Y. George, "A Coarse-to-Fine 3D U-Net Network for Semantic Segmentation of Kidney CT Scans," in Lecture notes in computer science, 2022, pp. 137-142. doi: 10.1007/978-3-030-98385-7_18.
- [41] Z. Zhao, H. Chen, J. Li, and L. Wang, "Boundary Attention U-Net for Kidney and Kidney Tumor Segmentation," 2022 44th Annual International Conference of the IEEE Engineering in Medicine & Biology Society (EMBC), vol. 2022, pp. 1540-1543, Jul. 2022, doi: 10.1109/embc48229.2022.9871443.

- [42] GeeksforGeeks, "Unet++ Architecture Explained," GeeksforGeeks, Jul. 31, 2023. <https://www.geeksforgeeks.org/machine-learning/unet-architecture-explained/>
- [43] N. Tomar, "Attention UNET and its Implementation in TensorFlow," Idiot Developer, Feb. 26, 2023. <https://idiotdeveloper.com/attention-unet-and-its-implementation-in-tensorflow/>
- [44] W. Zhao, D. Jiang, J. P. Queralta, and T. Westerlund, "MSS U-Net: 3D segmentation of kidneys and tumors from CT images with a multi-scale supervised U-Net," *Informatics in Medicine Unlocked*, vol. 19, p. 100357, Jan. 2020, doi: 10.1016/j.imu.2020.100357
- [45] K.-H. Uhm et al., "Exploring 3D U-Net Training Configurations and Post-processing Strategies for the MICCAI 2023 Kidney and Tumor Segmentation Challenge," in *Lecture notes in computer science*, 2024, pp. 8-13. doi: 10.1007/978-3-031-54806-2_2.
- [46] J. Chen et al., "TransUNET: Transformers make strong encoders for medical image segmentation," *arXiv.org*, Feb. 08, 2021. <https://arxiv.org/abs/2102.04306>
- [47] H. Cao et al., "Swin-Unet: Unet-like Pure Transformer for Medical Image Segmentation," *arXiv.org*, May 12, 2021. <https://arxiv.org/abs/2105.05537>
- [48] W. Hu, S. Yang, W. Guo, N. Xiao, X. Yang, and X. Ren, "STC-UNet: renal tumor segmentation based on enhanced feature extraction at different network levels," *BMC Medical Imaging*, vol. 24, no. 1, p. 179, Jul. 2024, doi: 10.1186/s12880-024-01359-5.
- [49] H. Qi, Z. Wang, X. Qi, Y. Shi, and T. Xie, "Ultrasound image segmentation of renal tumors based on UNet++ with fusion of multiscale residuals and dual attention," *Physics in Medicine and Biology*, vol. 69, no. 7, p. 075002, Feb. 2024, doi: 10.1088/1361-6560/ad2d7f.

- [50] “StructSeg2019 - Grand Challenge,” [grand-challenge.org](https://structseg2019.grand-challenge.org/Evaluation/).
<https://structseg2019.grand-challenge.org/Evaluation/>
- [51] F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, "nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation," *Nature Methods*, vol. 18, no. 2, pp. 203-211, Dec. 2020, doi: <https://doi.org/10.1038/s41592-020-01008-z>.
- [52] F. Isensee et al., "NNU-Net revisited: A call for rigorous validation in 3D medical image segmentation," *arXiv.org*, Apr. 15, 2024. <https://arxiv.org/abs/2404.09556>
- [53] P. Gupta, “nnU-Net : no-new-UNet,” *Medium*, Sep. 29, 2020. <https://medium.com/miccai-educational-initiative/nnu-net-the-no-new-unet-for-automatic-segmentation-8d655f3f6d2a>
- [54] J. Becktepe, L. Hennig, S. Oeltze-Jafra, and M. Lindauer, "Auto-NNU-Net: towards Automated medical image segmentation," *arXiv (Cornell University)*, May 2025, doi: [10.48550/arxiv.2505.16561](https://arxiv.org/abs/2505.16561).
- [55] P. M. Full, F. Isensee, P. F. Jager, and K. Maier-Hein, "Studying Robustness of Semantic Segmentation under Domain Shift in cardiac MRI," *arXiv.org*, Nov. 15, 2020. <https://arxiv.org/abs/2011.07592>
- [56] Y. Wang, Y. Dai, J. Zhang, and J. Yin, "Cascaded NNU-Net for kidney and kidney tumor segmentation," in *Lecture notes in computer science*, 2024, pp. 114-119. doi: [10.1007/978-3-031-54806-2_16](https://doi.org/10.1007/978-3-031-54806-2_16).
- [57] L. C. Moore et al., “Neural network dose prediction for cervical brachytherapy: Overcoming data scarcity for applicator-specific models,” *Medical Physics*, vol. 51, no. 7, pp. 4591–4606, May 2024, doi: <https://doi.org/10.1002/mp.17230>.

- [58] Z. Zhao, H. Chen, and L. Wang, "A Coarse-to-Fine Framework for the 2021 Kidney and Kidney Tumor Segmentation Challenge," in Lecture notes in computer science, 2022, pp. 53-58. doi: 10.1007/978-3-030-98385-7_8.
- [59] N. Heller et al., "The KiTS21 Challenge: Automatic segmentation of kidneys, renal tumors, and renal cysts in corticomedullary-phase CT," arXiv.org, Jul. 05, 2023. <https://arxiv.org/abs/2307.01984>
- [60] A. Soleimany, A. Karimian, M. Noroozi, and J. Majidpour, "Benchmarking nnU-Net vs. Custom 3D U-Net for Kidney Tumor Segmentation: A Controlled Study on KiTS19 Dataset," in 2025 32nd National and 10th International Iranian Conference on Biomedical Engineering (ICBME), 2025, doi: 10.1109/icbme68496.2025.11392310.
- [61] J. Kachele, M. Zenk, M. Rokuss, C. Ulrich, T. Wald, and K. H. Maier-Hein, "Enhanced NNU-Net architectures for automated MRI segmentation of head and neck tumors in adaptive radiation therapy," Lecture Notes in Computer Science, vol. 15273, pp. 50-64, Jan. 2025, doi: 10.1007/978-3-031-83274-1_3.
- [62] Y. Kirchhoff et al., "Optimizing nnU-Net with OpenVINO for Fast CPU Inference in Abdominal Organ Segmentation," in Lecture notes in computer science, 2025, pp. 96-105. doi: 10.1007/978-3-031-96202-8_9.
- [63] S. Doostinia et al., "Investigating the Self-optimizing nnU-NetV2 for Kidney Tumor Segmentation: Application to the KiTS23 Dataset," in 2025 32nd National and 10th International Iranian Conference on Biomedical Engineering (ICBME), 2025, doi: 10.1109/icbme68496.2025.11392506.
- [64] K. Yildirim, P. G. Bozdog, M. Talo, O. Yildirim, M. Karabatak, and U. Rajendra. Acharya, "Deep Learning Model for Automated Kidney Stone Detection Using Coronal

CT Images," *Computers in Biology and Medicine*, vol. 135, p. 104569, Aug. 2021, doi: <https://doi.org/10.1016/j.combiomed.2021.104569>.

[65] S. de Boer et al., "Robust Kidney Abnormality Segmentation: A Validation Study of an AI-Based Framework," *arXiv preprint arXiv:2505.07573*, 2025.

